

cleaning-dataset

June 9, 2024

```
[ ]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import folium
```

Celem jest predykcja ceny nieruchomości.

Zbiór danych zawiera 4802 instancje i kolumny określające:

BROKERTITLE: Title of the broker TYPE: Type of the house PRICE: Price of the house BEDS: Number of bedrooms BATH: Number of bathrooms PROPERTYSQFT: Square footage of the property ADDRESS: Full address of the house STATE: State of the house MAIN_ADDRESS: Main address information ADMINISTRATIVE_AREA_LEVEL_2: Administrative area level 2 information LOCALITY: Locality information SUBLOCALITY: Sublocality information STREET_NAME: Street name LONG_NAME: Long name FORMATTED_ADDRESS: Formatted address LATITUDE: Latitude coordinate of the house LONGITUDE: Longitude coordinate of the house

```
[ ]: import pandas as pd
url = 'NY-House-Dataset.csv'
data = pd.read_csv(url, sep= ';')
```

```
[ ]: data
```

```
[ ]:
      BROKERTITLE      TYPE \
0    Brokered by Douglas Elliman -111 Fifth Ave    Condo for sale
1              Brokered by Serhant    Condo for sale
2    Brokered by Sowae Corp    House for sale
3    Brokered by COMPASS    Condo for sale
4    Brokered by Sotheby's International Realty - E...    Townhouse for sale
...
4796    Brokered by COMPASS    Co-op for sale
4797    Brokered by Mjr Real Estate Llc    Co-op for sale
4798    Brokered by Douglas Elliman - 575 Madison Ave    Co-op for sale
4799    Brokered by E Realty International Corp    Condo for sale
4800    Brokered by Nyc Realty Brokers Llc    Co-op for sale

      PRICE  BEDS      BATH  PROPERTYSQFT \
0    315000     2    2.000000    1400.000000
```

1	195000000	7	10.000000	17545.000000
2	260000	4	2.000000	2015.000000
3	69000	3	1.000000	445.000000
4	55000000	7	2.373861	14175.000000
...
4796	599000	1	1.000000	2184.207862
4797	245000	1	1.000000	2184.207862
4798	1275000	1	1.000000	2184.207862
4799	598125	2	1.000000	655.000000
4800	349000	1	1.000000	750.000000

	ADDRESS \
0	2 E 55th St Unit 803
1	Central Park Tower Penthouse-217 W 57th New Yo...
2	620 Sinclair Ave
3	2 E 55th St Unit 908W33
4	5 E 64th St
...	...
4796	222 E 80th St Apt 3A
4797	97-40 62 Dr Unit Lg
4798	427 W 21st St Unit Garden
4799	91-23 Corona Ave Unit 4G
4800	460 Neptune Ave Apt 140

	STATE \
0	New York, NY 10022
1	New York, NY 10019
2	Staten Island, NY 10312
3	Manhattan, NY 10022
4	New York, NY 10065
...	...
4796	Manhattan, NY 10075
4797	Rego Park, NY 11374
4798	New York, NY 10011
4799	Elmhurst, NY 11373
4800	Brooklyn, NY 11224

	MAIN_ADDRESS \
0	2 E 55th St Unit 803New York, NY 10022
1	Central Park Tower Penthouse-217 W 57th New Yo...
2	620 Sinclair AveStaten Island, NY 10312
3	2 E 55th St Unit 908W33Manhattan, NY 10022
4	5 E 64th StNew York, NY 10065
...	...
4796	222 E 80th St Apt 3AManhattan, NY 10075
4797	97-40 62 Dr Unit LgRego Park, NY 11374
4798	427 W 21st St Unit GardenNew York, NY 10011

4799 91-23 Corona Ave Unit 4GElmhurst, NY 11373
 4800 460 Neptune Ave Apt 140Brooklyn, NY 11224

	ADMINISTRATIVE_AREA_LEVEL_2	LOCALITY	SUBLOCALITY \
0	New York County	New York	Manhattan
1	United States	New York	New York County
2	United States	New York	Richmond County
3	United States	New York	New York County
4	United States	New York	New York County
...
4796	New York	New York County	New York
4797	United States	New York	Queens County
4798	United States	New York	New York County
4799	New York	Queens County	Queens
4800	New York	Kings County	Brooklyn

	STREET_NAME	LONG_NAME \
0	East 55th Street	Regis Residence
1	New York	West 57th Street
2	Staten Island	Sinclair Avenue
3	New York	East 55th Street
4	New York	East 64th Street
...
4796	Manhattan	222
4797	Queens	62nd Drive
4798	New York	West 21st Street
4799	Flushing	91-23
4800	Coney Island	460

	FORMATTED_ADDRESS	LATITUDE	LONGITUDE
0	Regis Residence, 2 E 55th St #803, New York, N...	40.761255	-73.974483
1	217 W 57th St, New York, NY 10019, USA	40.766393	-73.980991
2	620 Sinclair Ave, Staten Island, NY 10312, USA	40.541805	-74.196109
3	2 E 55th St, New York, NY 10022, USA	40.761398	-73.974613
4	5 E 64th St, New York, NY 10065, USA	40.767224	-73.969856
...
4796	222 E 80th St #3a, New York, NY 10075, USA	40.774350	-73.955879
4797	97-40 62nd Dr, Rego Park, NY 11374, USA	40.732538	-73.860152
4798	427 W 21st St, New York, NY 10011, USA	40.745882	-74.003398
4799	91-23 Corona Ave. #4b, Flushing, NY 11373, USA	40.742770	-73.872752
4800	460 Neptune Ave #14a, Brooklyn, NY 11224, USA	40.579147	-73.970949

[4801 rows x 17 columns]

Podział kolumn na dwie kategorie: dane numeryczne i dane katagoryczne.

```
[ ]: numeric_columns = {'PRICE', 'BEDS', 'BATH', 'PROPERTYSQFT', 'LATITUDE', 'LONGITUDE'}
category_columns = {'TYPE', 'BROKERTITLE', 'STATE', 'LONG_NAME', 'LOCALITY', 'ADDRESS', 'ADMINISTRATIVE_AREA_LEVEL_2', 'SUBLOCALITY', 'MAIN_ADDRESS', 'FORMATTED_ADDRESS', 'STREET_NAME'}
```

Zmianienie liter na małe w celu ujednolicenia danych.

```
[ ]: for column in category_columns:
    data[column] = data[column].str.lower()
```

Sprawdzenie unikalnych wartości oraz charakterystyki zbioru danych.

```
[ ]: data[list(category_columns)].nunique()
```

```
[ ]: STATE                308
FORMATTED_ADDRESS        4550
LONG_NAME                2731
ADDRESS                  4582
LOCALITY                 11
BROKERTITLE              1011
SUBLOCALITY              21
MAIN_ADDRESS             4583
ADMINISTRATIVE_AREA_LEVEL_2 29
TYPE                     13
STREET_NAME              174
dtype: int64
```

```
[ ]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4801 entries, 0 to 4800
Data columns (total 17 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   BROKERTITLE                          4801 non-null   object
1   TYPE                                4801 non-null   object
2   PRICE                               4801 non-null   int64
3   BEDS                                4801 non-null   int64
4   BATH                                4801 non-null   float64
5   PROPERTYSQFT                        4801 non-null   float64
6   ADDRESS                             4801 non-null   object
7   STATE                               4801 non-null   object
8   MAIN_ADDRESS                        4801 non-null   object
9   ADMINISTRATIVE_AREA_LEVEL_2         4801 non-null   object
10  LOCALITY                            4801 non-null   object
11  SUBLOCALITY                         4801 non-null   object
12  STREET_NAME                         4801 non-null   object
```

```

13 LONG_NAME          4801 non-null object
14 FORMATTED_ADDRESS  4801 non-null object
15 LATITUDE           4801 non-null float64
16 LONGITUDE          4801 non-null float64
dtypes: float64(4), int64(2), object(11)
memory usage: 637.8+ KB

```

```
[ ]: data.describe().T
```

```

[ ]:
      count      mean      std      min      25% \
PRICE      4801.0  2.356940e+06  3.135525e+07  2494.000000  499000.000000
BEDS       4801.0  3.356801e+00  2.602315e+00    1.000000    2.000000
BATH       4801.0  2.373861e+00  1.946962e+00    0.000000    1.000000
PROPERTYSQFT 4801.0  2.184208e+03  2.377141e+03   230.000000   1200.000000
LATITUDE    4801.0  4.071423e+01  8.767557e-02   40.499546   40.639375
LONGITUDE   4801.0 -7.394160e+01  1.010825e-01  -74.253033  -73.987143

      50%      75%      max
PRICE      825000.000000  1.495000e+06  2.147484e+09
BEDS           3.000000  4.000000e+00  5.000000e+01
BATH           2.000000  3.000000e+00  5.000000e+01
PROPERTYSQFT  2184.207862  2.184208e+03  6.553500e+04
LATITUDE       40.726749  4.077192e+01  4.091273e+01
LONGITUDE     -73.949189 -7.387064e+01 -7.370245e+01

```

```
[ ]: data.head()
```

```

[ ]:
      BROKERTITLE      TYPE \
0  brokered by douglas elliman -111 fifth ave  condo for sale
1  brokered by serhant  condo for sale
2  brokered by sowae corp  house for sale
3  brokered by compass  condo for sale
4  brokered by sotheby's international realty - e...  townhouse for sale

      PRICE  BEDS  BATH  PROPERTYSQFT \
0    315000    2    2.000000    1400.0
1  195000000    7   10.000000   17545.0
2    260000    4    2.000000    2015.0
3    69000    3    1.000000     445.0
4  55000000    7    2.373861   14175.0

      ADDRESS      STATE \
0    2 e 55th st unit 803  new york, ny 10022
1  central park tower penthouse-217 w 57th new yo...  new york, ny 10019
2    620 sinclair ave  staten island, ny 10312
3    2 e 55th st unit 908w33  manhattan, ny 10022
4    5 e 64th st  new york, ny 10065

```

	MAIN_ADDRESS \			
0	2 e 55th st unit 803new york, ny 10022			
1	central park tower penthouse-217 w 57th new yo...			
2	620 sinclair avestaten island, ny 10312			
3	2 e 55th st unit 908w33manhattan, ny 10022			
4	5 e 64th stnew york, ny 10065			

	ADMINISTRATIVE_AREA_LEVEL_2	LOCALITY	SUBLOCALITY	STREET_NAME \
0	new york county	new york	manhattan	east 55th street
1	united states	new york	new york county	new york
2	united states	new york	richmond county	staten island
3	united states	new york	new york county	new york
4	united states	new york	new york county	new york

	LONG_NAME	FORMATTED_ADDRESS \
0	regis residence	regis residence, 2 e 55th st #803, new york, n...
1	west 57th street	217 w 57th st, new york, ny 10019, usa
2	sinclair avenue	620 sinclair ave, staten island, ny 10312, usa
3	east 55th street	2 e 55th st, new york, ny 10022, usa
4	east 64th street	5 e 64th st, new york, ny 10065, usa

	LATITUDE	LONGITUDE
0	40.761255	-73.974483
1	40.766393	-73.980991
2	40.541805	-74.196109
3	40.761398	-73.974613
4	40.767224	-73.969856

Sprawdzenie braków danych.

```
[ ]: data.isna().sum()
```

```
[ ]: BROKERTITLE      0
      TYPE            0
      PRICE           0
      BEDS            0
      BATH            0
      PROPERTYSQFT    0
      ADDRESS         0
      STATE           0
      MAIN_ADDRESS    0
      ADMINISTRATIVE_AREA_LEVEL_2  0
      LOCALITY        0
      SUBLOCALITY     0
      STREET_NAME     0
      LONG_NAME       0
```

```

FORMATTED_ADDRESS      0
LATITUDE               0
LONGITUDE              0
dtype: int64

```

Sprawdzenie ilości zduplikowanych wierszy i ich usunięcie.

```

[ ]: print('Duplicated rows: ', data.duplicated().sum())
data.drop_duplicates(inplace=True)

```

Duplicated rows: 214

Usunięcie w kolumnie „BROKERTITLE” ciągów znaków ‘llc’, ‘inc’, zamiana wartości ‘rlty’ na ‘realty’ i zapisanie wyników tych zmian w nowej kolumnie „Broker”.

BROKERTITLE

```

[ ]: data['BROKER'] = data['BROKERTITLE'].str.replace('llc','')
data['BROKER'] = data['BROKER'].str.replace('inc','')
data['BROKER'] = data['BROKER'].str.replace('rlty','realty')
data['BROKER'] = data['BROKER'].str.replace('.', '')
def split_by_delimeter(value, separator):
    result = value.split(separator)[0] if separator in value else value
    result = result.strip()
    return result

data['BROKER'] = data['BROKER'].apply(lambda x: split_by_delimeter(x, ' -'))

```

TYPE

```

[ ]: data.loc[data['TYPE'] == 'land for sale', ['BATH', 'BEDS']]

```

W kolumnie „TYPE” ciąg znaków “condop” zmieniony na “condo”.

Stworzenie nowej kolumny „ANNOUNCEMENT_TYPE” która przypisuje ogłoszeniom odpowiednie kategorie na podstawie ich typu (‘TYPE’) przy użyciu wcześniej zdefiniowanego słownika.

Usuujemy również ciąg znaków „for sale” z wartości w kolumnie TYPE.

```

[ ]: #Zamieniamy liczbe łazienek i sypialni na zero dla ogłoszeń, które są dla
    ↪sprzedaży działki
data.loc[data['TYPE'] == 'land for sale', ['BATH', 'BEDS']] = 0
data["TYPE"] = data["TYPE"].str.replace('condop', 'condo')

data['ANNOUNCEMENT_TYPE'] = data['TYPE'].replace({
    'condo for sale': 'apartment',
    'townhouse for sale': 'home',
    'house for sale': 'home',
    'multi-family home for sale': 'home',
    'co-op for sale': 'co-op',
})

```

```

        'mobile house for sale': 'home',
        'land for sale': 'land',
        'foreclosure': 'other',
        'contingent': 'other',
        'pending': 'other',
        'coming soon': 'other',
        'for sale': 'other'})
data["TYPE"] = data["TYPE"].str.replace(" for sale", "")

```

PRICE - W kolumnie „PRICE” usuwamy wiersze, dla których cena jest większa niż 100 000 000 lub mniejsza niż 10 000 i zmienienie typ danych na liczby zmiennoprzecinkowe.

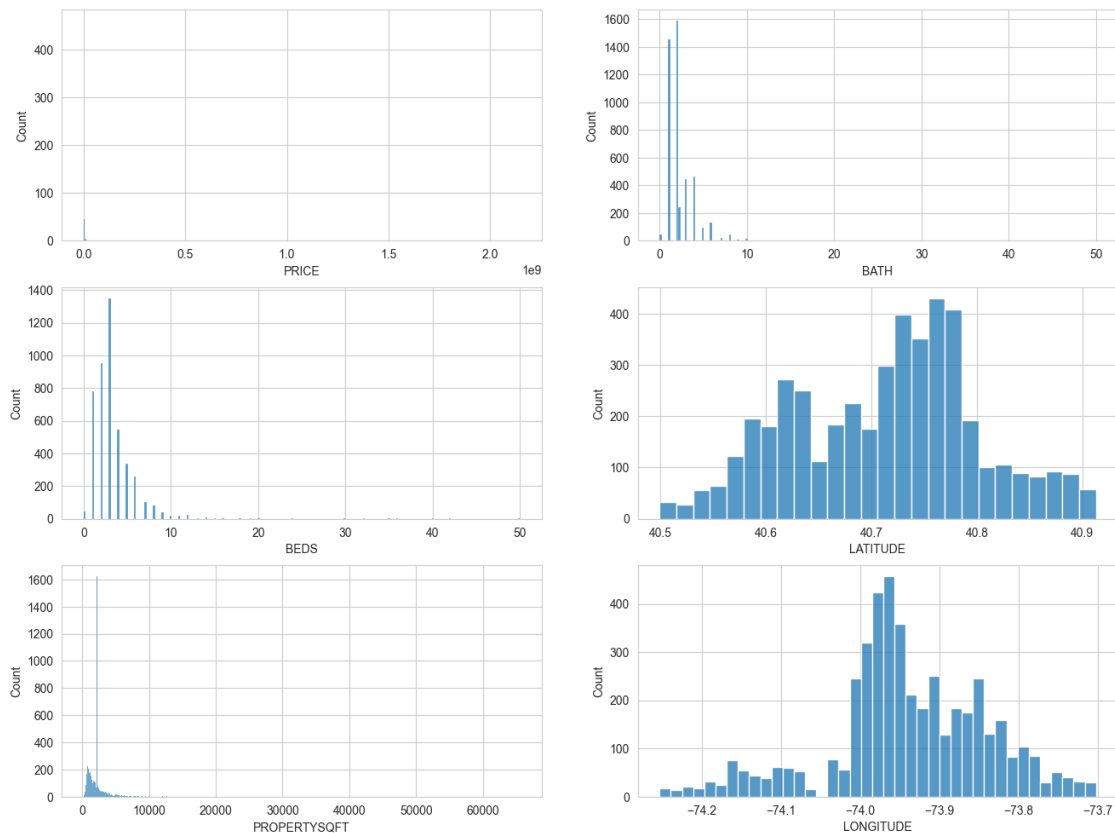
```

[ ]: fig, axes = plt.subplots(3,2, figsize=(16, 12))
    axes = axes.flatten()

    i=0
    for column in numeric_columns:
        sns.histplot(x=data[column], ax=axes[i])
        i=i+1

    plt.show()

```




```
[ ]: data.drop(data.loc[(data["PRICE"] > 100000000) | (data["PRICE"]<10000)].index,
    inplace= True)
data['PRICE'] = data['PRICE'].astype(float)
```

BATH, BEDS

```
[ ]: data['BATH'].value_counts()
```

```
[ ]: data['BEDS'].value_counts()
```

Utworzenie nowej kolumny „PROPERTYSQFT1000”, w której wartości są obliczane przez podzielenie wartości w kolumnie ‘PROPERTYSQFT’ przez 1000, zaokrąglając wynik w dół do liczby całkowitej , w celu znalezienia wartości odstających.

Zostały pozostawione wiersze z wartościami dla których metraż nieruchomości podzielony przez 1000 mieści się w zakresie od 0 do 8.

Pozostawiono wiersze z kolumny „BATH”, w których liczba łazienek mieści się w zakresie od 0 do 10, ze względu na małą ilość wystąpień wartości spoza tego przedziału.

Pozostawiono wiersze z kolumny „BEDS” w których liczba sypialni mieści się w zakresie od 0 do 12, ze względu na małą ilość wystąpień wartości spoza tego przedziału.

```
[ ]: data['PROPERTYSQFT1000'] = data['PROPERTYSQFT'].apply(lambda x: x//1000)
data['PROPERTYSQFT1000'].value_counts()
```

```
[ ]: data = data[(data['BATH'] >= 0) & (data['BATH'] <= 10)]
data = data[(data['BEDS'] >= 0) & (data['BEDS'] <= 12)]
data['BATH'] = data['BATH'].apply(lambda x: x if float.is_integer(x) else
    int(x))
data = data[(data['PROPERTYSQFT1000'] >= 0) & (data['PROPERTYSQFT1000'] <= 8)]

data['PROPERTYSQFT'] = data['PROPERTYSQFT'].round(2)
```

Cleaning adresses

```
[ ]: data_api = data.copy()
```

Usunięcie kolumny „MAIN ADDRESS”, ponieważ jest połączeniem kolumn „ADDRESS” i „STATE”.

```
[ ]: (data["ADDRESS"] + data["STATE"] == data["MAIN_ADDRESS"]).unique()
data.drop("MAIN_ADDRESS", axis=1, inplace=True)
```

Utworzenie słownika „dictLocality”, który mapuje nazwy hrabstw na nazwy dzielnic w Nowym Jorku.

Stworzenie listy „listOfBoroughs” zawierającej nazwy dzielnic (wartości powyżej wspomnianego słownika).

```
[ ]: dictLocality = {"new york county": "manhattan", "kings county":
    ↪ "brooklyn", "bronx county": "the bronx", "richmond county": "staten island",
    ↪ "queens county": "queens"}
listOfBoroughs = list(dictLocality.values())
```

Zamieniamy wartości występujące jako klucze w słowniku na odpowiadające im wartości we wszystkich kolumnach.

```
[ ]: for c in list(data.columns):
    for i in list(dictLocality.keys()):
        data[c] = data[c].astype(str).str.replace(i, dictLocality[i])
```

```
[ ]: # Słownik mapujący skrócone i niepoprawne nazwy stanów i dzielnic na ich
    ↪ poprawne nazwy
state_replacements = {
    'nyc': 'new york',
    'ny': 'new york',
    'new yorkc': 'new york',
    'new york city': 'new york',
    r'kew gardens hill(?:s)': 'kew garden hills',
    r'kew gardens(?: hills)': 'kew garden hills',
    'kew gardens hills': 'kew garden hills',
    'kew gardens hillss': 'kew garden hills',
    r'(?<!\bthe\s)bronx new york': 'the bronx',
    'queens village': 'queens',
    r'(?<!\bthe\s)bronx ny': 'the bronx',
    'brooklyn heights': 'brooklyn',
    r'\b(?<!\bthe\s)(?<!\bwest\s)bronx\b': 'the bronx'
}

for col in ["ADDRESS", "ADMINISTRATIVE_AREA_LEVEL_2", "STATE", "LOCALITY",
    ↪ "SUBLOCALITY", "STREET_NAME", "LONG_NAME", "FORMATTED_ADDRESS"]:
    for i in list(state_replacements.keys()):
        data[col] = data[col].str.replace(i, state_replacements[i], regex=True)
```

```
[ ]: print(data["LOCALITY"].unique())
# Przypisanie wartości z kolumny "LOCALITY" do kolumny "BOROUGH", jeśli wartość
    ↪ znajduje się w liście listOfBoroughs.
data.loc[data["LOCALITY"].isin(listOfBoroughs), "BOROUGH"] = data['LOCALITY']
print(data["BOROUGH"].unique())
# Jeżeli wartość w kolumnie LOCALITY to flatbush to przepiszujemy tą wartość do
    ↪ kolumny NEIGHBOURHOOD
data.loc[data["LOCALITY"] == "flatbush", "NEIGHBOURHOOD"] = data['LOCALITY']
```

```
['new york' 'manhattan' 'the bronx' 'brooklyn' 'queens' 'staten island'
 'united states' 'flatbush']
```

```
[nan 'manhattan' 'the bronx' 'brooklyn' 'queens' 'staten island']
```

```
[ ]: print(data["SUBLOCALITY"].unique())
# Przepisanie wartości z SUBLOCALITY do kolumny BOROUGH, jeśli wartość jest
  ↳ zawarta w listOfBoroughs oraz komórka w kolumnie BOROUGH jest pusta.
data.loc[(data["SUBLOCALITY"].isin(listOfBoroughs)) & (data["BOROUGH"].isna()) ,
  ↳ "BOROUGH"] = data["SUBLOCALITY"]

# Jeżeli kolumna SUBLOCALITY zawiera "snyder avenue" to przepisujemy tą wartość
  ↳ do kolumny STREET.
data.loc[data["SUBLOCALITY"] == "snyder avenue" , "STREET"] =
  ↳ data["SUBLOCALITY"]

# Jeśli wartość w kolumnie SUBLOCALITY nie znajduje się w listOfBoroughs, nie
  ↳ jest równa "new york" ani "snyder avenue" to przypisujemy tą wartość do
  ↳ odpowiadającej komórki w kolumnie NEIGHBOURHOOD.
data.loc[(data["SUBLOCALITY"].isin(list(filter(lambda x: x not in
  ↳ listOfBoroughs, list(data["SUBLOCALITY"])))))) & (~(data["SUBLOCALITY"].
  ↳ isin(["new york", "snyder avenue"]))) , "NEIGHBOURHOOD"] =
  ↳ data["SUBLOCALITY"]
```

```
['manhattan' 'staten island' 'brooklyn' 'new york' 'east bronx'
 'the bronx' 'queens' 'coney island' 'jackson heights' 'riverdale'
 'rego park' 'fort hamilton' 'flushing' 'dumbo' 'snew yorkder avenue']
```

```
[ ]: data
```

```
[ ]:
      BROKERTITLE  TYPE  PRICE BEDS \
0    brokered by douglas elliman -111 fifth ave  condo  315000.0  2
2                brokered by sowae corp  house  260000.0  4
3                brokered by compass  condo   69000.0  3
5                brokered by sowae corp  house  690000.0  5
6    brokered by douglas elliman - 575 madison ave  condo  899500.0  2
...
4796                brokered by compass  co-op   599000.0  1
4797                brokered by mjr real estate llc  co-op   245000.0  1
4798    brokered by douglas elliman - 575 madison ave  co-op  1275000.0  1
4799                brokered by e realty international corp  condo   598125.0  2
4800                brokered by nyc realty brokers llc  co-op   349000.0  1

      BATH  PROPERTYSQFT  ADDRESS \
0    2.0    1400.0    2 e 55th st unit 803
2    2.0    2015.0    620 sinclair ave
3    1.0    445.0    2 e 55th st unit 908w33
5    2.0    4004.0    584 park pl
6    2.0    2184.21    157 w 126th st unit 1b
...    ...    ...    ...
```

4796	1.0	2184.21	222 e 80th st apt 3a
4797	1.0	2184.21	97-40 62 dr unit lg
4798	1.0	2184.21	427 w 21st st unit garden
4799	1.0	655.0	91-23 corona ave unit 4g
4800	1.0	750.0	460 neptune ave apt 14o

		STATE ADMINISTRATIVE_AREA_LEVEL_2	LOCALITY \
0	new york, new york 10022	manhattan	new york
2	staten island, new york 10312	united states	new york
3	manhattan, new york 10022	united states	new york
5	brooklyn, new york 11238	united states	new york
6	new york, new york 10027	new york	manhattan
...
4796	manhattan, new york 10075	new york	manhattan
4797	rego park, new york 11374	united states	new york
4798	new york, new york 10011	united states	new york
4799	elmhurst, new york 11373	new york	queens
4800	brooklyn, new york 11224	new york	brooklyn

	LONG_NAME \
0	... regis residence
2	... sinclair avenue
3	... east 55th street
5	... park place
6	... 157
...	...
4796	... 222
4797	... 62nd drive
4798	... west 21st street
4799	... 91-23
4800	... 460

	FORMATTED_ADDRESS	LATITUDE \
0	regis residence, 2 e 55th st #803, new york, n...	40.761255
2	620 sinclair ave, staten island, new york 1031...	40.5418051
3	2 e 55th st, new york, new york 10022, usa	40.7613979
5	584 park pl, brooklyn, new york 11238, usa	40.6743632
6	157 w 126th st #1b, new york, new york 10027, usa	40.809448
...
4796	222 e 80th st #3a, new york, new york 10075, usa	40.77435
4797	97-40 62nd dr, rego park, new york 11374, usa	40.7325379
4798	427 w 21st st, new york, new york 10011, usa	40.7458817
4799	91-23 corona ave. #4b, flushing, new york 1137...	40.7427705
4800	460 neptune ave #14a, brooklyn, new york 11224...	40.579147

	LONGITUDE	BROKER ANNOUNCEMENT_TYPE \
0	-73.9744834	brokered by douglas elliman apartment

2	-74.1961086	brokered by sowae corp	home
3	-73.9746128	brokered by compass	apartment
5	-73.9587248	brokered by sowae corp	home
6	-73.946777	brokered by douglas elliman	apartment
...
4796	-73.955879	brokered by compass	co-op
4797	-73.8601516	brokered by mjr real estate	co-op
4798	-74.0033976	brokered by douglas elliman	co-op
4799	-73.8727516	brokered by e realty international corp	apartment
4800	-73.9709488	brokered by nyc realty brokers	co-op

	PROPERTY	SQFT	1000	BOROUGH	NEIGHBOURHOOD	STREET
0		1.0		manhattan	NaN	NaN
2		2.0		staten island	NaN	NaN
3		0.0		manhattan	NaN	NaN
5		4.0		brooklyn	NaN	NaN
6		2.0		manhattan	NaN	NaN
...	
4796		2.0		manhattan	NaN	NaN
4797		2.0		queens	NaN	NaN
4798		2.0		manhattan	NaN	NaN
4799		0.0		queens	NaN	NaN
4800		0.0		brooklyn	NaN	NaN

[4507 rows x 22 columns]

```
[ ]: print(data["ADMINISTRATIVE_AREA_LEVEL_2"].unique())
#Jeśli wartość z kolumny ADMINISTRATIVE_AREA_LEVEL_2 znajduje się w
↳listOfBoroughs, a kolumna "BOROUGH" jest pusta to uzupełniamy kolumnę
↳Borough sprawdzoną wartością.
data.loc[(data["ADMINISTRATIVE_AREA_LEVEL_2"].isin(listOfBoroughs))
↳&(data["BOROUGH"].isna()) , "BOROUGH"] = data["ADMINISTRATIVE_AREA_LEVEL_2"]

# Uzupełnienie kolumny POSTCODE wartościami z kolumny
↳ADMINISTRATIVE_AREA_LEVEL_2, jeśli nie są zawarte w listOfBoroughs, nie są
↳równe "new york" ani "united states".
data.loc[(data["ADMINISTRATIVE_AREA_LEVEL_2"].isin(list(filter(lambda x: x not
↳in listOfBoroughs, list(data["ADMINISTRATIVE_AREA_LEVEL_2"])))) &
↳(data["ADMINISTRATIVE_AREA_LEVEL_2"] != "new york") &
↳(data["ADMINISTRATIVE_AREA_LEVEL_2"] != "united states") , "POSTCODE"] =
↳data["ADMINISTRATIVE_AREA_LEVEL_2"]
```

```
['manhattan' 'united states' 'new york' 'the bronx' '11214' '10301'
'10309' '10303' '11234' '11414' '10310' '10003' '11417' '10304'
'brooklyn' '10463' 'queens' '10017' '10306' '10471' '11229' '10312'
'11412' '10465' '10002' '10466' '11237' '11218']
```

```
[ ]: data["STREET_NAME"].unique()
```

```
[ ]: array(['east 55th street', 'staten island', 'new york', 'brooklyn',  
          'manhattan', 'morrison avenue', 'midwood', 'concourse village',  
          'flushing', 'elmhurst', 'annadale', 'queens', 'fort hamilton',  
          'north riverdale', 'rego park', 'forest hills', 'the bronx',  
          'dongan hills', 'jackson heights', 'clifton', 'mariners harbor',  
          'dyker heights', 'williamsburg', 'concourse', 'mid island',  
          'centre street', 'cobble hill', 'park slope', 'brighton beach',  
          'flatbush', 'prospect heights', 'woodhaven', 'bedford-stuyvesant',  
          'jamaica', 'spuyten duyvil', 'bay ridge', 'shore acres', 'bayside',  
          'glen oaks', 'fresh meadows', 'highbridge', 'sheepshead bay',  
          'rector place', 'kew garden hills', 'bushwick', 'hudson hill',  
          'rosedale', 'east bronx', 'parkchester', 'borough park',  
          'little haiti', 'canarsie', 'kensington', 'east 110th street',  
          'east new york', 'pelham bay', 'howard beach', 'downtown brooklyn',  
          'city island', 'fieldston', 'westchester square', 'rockaway park',  
          'riverdale', 'bulls head', 'gerritsen beach', 'crown heights',  
          'new dorp', 'west bronx', 'ocean hill', 'gravesend',  
          'castleton corners', 'seagate', 'gowanus', 'windsor terrace',  
          'bath beach', 'norwood', 'whitestone', 'surf avenue',  
          'great kills', 'homecrest', 'long island city', 'woodside',  
          'maspeth', 'sunset park', 'douglaston', 'astoria', 'dumbo',  
          'new springville', 'corona', 'madison', 'bensonhurst',  
          'fordham manor', 'greenpoint', 'coney island', 'mill basin',  
          'carroll gardens', 'old fulton street', 'ozone park',  
          'east 74th street', 'floral park', '35th avenue', 'far rockaway',  
          'beechhurst', 'henry hudson parkway', 'rosebank', 'kingsbridge',  
          'bergen beach', 'ridgewood', 'oakwood', 'clinton hill',  
          'richmond hill', 'auburndale', 'southside', 'little caribbean',  
          'peck slip', 'fort greene', 'boerum hill', 'foxburst', 'red hook',  
          'columbia street waterfront district', 'east flatbush',  
          'sunnew yorkside', 'flatlands', 'mapleton', '98th place',  
          'east 22nd street', 'manhattan beach', 'east elmhurst',  
          'oxford avenue', 'bay terrace', 'arverne', 'midland beach',  
          'vinegar hill', 'little neck', 'west brighton', 'east 96th street',  
          'clason point', '5th avenue', 'shore road', 'west 56th street',  
          'middle village', 'west 111th street', 'prospect lefferts gardens',  
          'woodstock', 'melrose', 'east 10th street', '139th street',  
          'park avenue', 'west 65th street', 'john street',  
          'east end avenue', 'brownsville', '3g', 'west 13th street',  
          'college point', 'central park west', 'east 88th street',  
          'allerton', 'morrissania', 'west 64th street', '61st street',  
          'cypress hills', '2501', '67th drive', 'todt hill',  
          'saunders street', 'mount eden'], dtype=object)
```

```
[ ]: len(data["STREET_NAME"].unique())
```

[]: 167

Utworzenie słownika z rozwinięciami skrótów i zamiana wartości z kolumny na wartości ze słownika. Nazwy ulicy zostają dodane do kolumny "streets", jeśli zawierają jedno ze słów kluczowych. Przypisanie wartości z kolumn do kolumn dla wierszy, w których wartość znajduje się na liście i jednocześnie wartość w kolumnie jest pusta np. kolumny "STREET_NAME" i "BOROUGH", "STREET_NAME" i "NEIGHBOURHOOD", "STREET_NAME" i "STREET". Uzupełnienie brakujących informacji w kolumnach "POSTCODE", "BOROUGH" i "NEIGHBOURHOOD". Przypisanie wartości do konkretnych kolumn.

```
[ ]: addresses_shortcut = {
    'st': 'street',
    'ave': 'avenue',
    'rd': 'road',
    'blvd': 'boulevard',
    'dr': 'drive',
    'pkwy': 'parkway',
    'ct': 'court',
    'ln': 'lane',
    'pl': 'place',
    'sq': 'square',
    'apt': '',
    'ste': '',
    'num': ''
}

for c in ["STREET_NAME", "ADDRESS", "STATE", "LONG_NAME", "FORMATTED_ADDRESS"]:
    for add in addresses_shortcut.keys():
        mask = data[c].str.contains(fr'\b{add}\b')
        data.loc[mask, c] = data.loc[mask, c].str.replace(add,
↳addresses_shortcut[add])

streets = []
keyWords = ["street", "parkway", "avenue", "drive", "road"]
for street_name in data["STREET_NAME"]:
    for keyword in keyWords:
        if keyword in street_name:
            streets.append(street_name)

streets
```

```
[ ]: # Jeśli wartość z kolumny STREET_NAME znajduje się w listOfBoroughs, a kolumna
↳BOROUGH jest pusta, uzupełniamy kolumnę BOROUGH tą wartością.
data.loc[(data["STREET_NAME"].isin(listOfBoroughs)) & (data["BOROUGH"].isna()),
↳"BOROUGH"] = data["STREET_NAME"]
```

```
# Uzupełnienie pustych wartości w kolumnie NEIGHBOURHOOD wartościami z kolumny
↳ STREET_NAME, jeśli nie jest zawarta w liście streets, nie jest równa "new
↳ york", nie należy do listOfBoroughs.
data.loc[(data["STREET_NAME"].isin(list(filter(lambda x: (x not in streets) &
↳ (x != "new york") & (x not in listOfBoroughs), list(data["STREET_NAME"])))) &
↳ (data["NEIGHBOURHOOD"].isna()), "NEIGHBOURHOOD"] = data["STREET_NAME"]
data.loc[(data["STREET_NAME"].isin(streets)) & (data["STREET"].isna()),
↳ "STREET"] = data["STREET_NAME"]
```

```
[ ]: print(data["STATE"].unique())
# Uzupełnienie pustych wartości w kolumnie POSTCODE wartościami z kolumny STATE.
data.loc[data["POSTCODE"].isna(), "POSTCODE"] = data["STATE"].str.slice(-5)

# Jeżeli BOROUGH jest pusta to weź wartość z kolumny STATE po przecinku - pod
↳ warunkiem że należy do listOfBoroughs
data.loc[(data["BOROUGH"].isna()) & (data["STATE"].str.split(", ").str.get(0).
↳ isin(listOfBoroughs)), "BOROUGH"] = data["STATE"].str.split(", ").str.get(0)

# Jeżeli NEIGHBOURHOOD jest pusta to weź wartość z kolumny STATE po przecinku -
↳ pod warunkiem że nie należy do listOfBoroughs ani nie jest równa "new york",
↳ "ny", "nyc"
data.loc[(data["NEIGHBOURHOOD"].isna()) & (~ (data["STATE"].str.split(", ").str.
↳ get(0).isin(listOfBoroughs))) & (~ (data["STATE"].str.split(", ").str.get(0).
↳ isin(["new york", "ny", "nyc"]))) , "NEIGHBOURHOOD"] = data["STATE"].str.
↳ split(", ").str.get(0)
```

```
[ ]: data.loc[data["POSTCODE"].isna(), "POSTCODE"] = data["FORMATTED_ADDRESS"].str.
↳ split(", ").str.get(-2).str.slice(-5)
data.loc[(data["BOROUGH"].isna()) & (data["FORMATTED_ADDRESS"].str.split(", ").
↳ str.get(-3).isin(listOfBoroughs)), "BOROUGH"] = data["FORMATTED_ADDRESS"].
↳ str.split(", ").str.get(-3)
data.loc[(data["NEIGHBOURHOOD"].isna()) & (~ (data["FORMATTED_ADDRESS"].str.
↳ split(", ").str.get(-3).isin(listOfBoroughs))) &
↳ (~ (data["FORMATTED_ADDRESS"].str.split(", ").str.get(-3).isin(["new
↳ york"]))) , "NEIGHBOURHOOD"] = data["FORMATTED_ADDRESS"].str.split(", ").str.
↳ get(-3)
```

```
[ ]: data.loc[data["FORMATTED_ADDRESS"].str.split(", ").str.get(-4).str.split().str.
↳ get(0).str.replace("-", "").str.contains('\d'), "HOUSE_NUMBER"] =
↳ data["FORMATTED_ADDRESS"].str.split(", ").str.get(-4).str.split().str.get(0)
data.loc[(data["STREET"].isna()) & (data["FORMATTED_ADDRESS"].str.split(", ").
↳ str.get(-4).str.split().str.get(0).str.replace("-", "").str.contains('\d')),
↳ "STREET"] = data["FORMATTED_ADDRESS"].str.split(", ").str.get(-4).str.
↳ split().str.slice(start=1).str.join(" ")
```



```
data.loc[(data["STREET"].isna()) & (~ (data["FORMATTED_ADDRESS"].str.split(", ").
↳str.get(-4).str.split().str.get(0).str.replace("-", "").str.contains('\d'))),
↳"STREET"] = data["FORMATTED_ADDRESS"].str.split(", ").str.get(-4)
```

```
<>:1: SyntaxWarning: invalid escape sequence '\d'
<>:2: SyntaxWarning: invalid escape sequence '\d'
<>:3: SyntaxWarning: invalid escape sequence '\d'
<>:1: SyntaxWarning: invalid escape sequence '\d'
<>:2: SyntaxWarning: invalid escape sequence '\d'
<>:3: SyntaxWarning: invalid escape sequence '\d'
/var/folders/h4/pdn3pcp16vxc0wjz5jwhfhm40000gn/T/ipykernel_48297/2895737151.py:1
: SyntaxWarning: invalid escape sequence '\d'
data.loc[data["FORMATTED_ADDRESS"].str.split(",
").str.get(-4).str.split().str.get(0).str.replace("-", "").str.contains('\d'),
"HOUSE_NUMBER"] = data["FORMATTED_ADDRESS"].str.split(",
").str.get(-4).str.split().str.get(0)
/var/folders/h4/pdn3pcp16vxc0wjz5jwhfhm40000gn/T/ipykernel_48297/2895737151.py:2
: SyntaxWarning: invalid escape sequence '\d'
data.loc[(data["STREET"].isna()) & (data["FORMATTED_ADDRESS"].str.split(",
").str.get(-4).str.split().str.get(0).str.replace("-", "").str.contains('\d')),
"STREET"] = data["FORMATTED_ADDRESS"].str.split(",
").str.get(-4).str.split().str.slice(start=1).str.join(" ")
/var/folders/h4/pdn3pcp16vxc0wjz5jwhfhm40000gn/T/ipykernel_48297/2895737151.py:3
: SyntaxWarning: invalid escape sequence '\d'
data.loc[(data["STREET"].isna()) & (~ (data["FORMATTED_ADDRESS"].str.split(",
").str.get(-4).str.split().str.get(0).str.replace("-", "").str.contains('\d'))),
"STREET"] = data["FORMATTED_ADDRESS"].str.split(", ").str.get(-4)
```

```
[ ]: data.loc[(data["LONG_NAME"].str.replace("-", "").str.isdigit()),
↳"HOUSE_NUMBER"] = data["LONG_NAME"]
data.loc[(~ (data["LONG_NAME"].str.replace("-", "").str.isdigit())) &
↳(data["STREET"].isna()) & (data["LONG_NAME"] != "parking lot"), "STREET"] =
↳data["LONG_NAME"]
```

```
[ ]: data.loc[(data["ADDRESS"].str.split(" ").str.get(0).str.replace("-", "").str.
↳isdigit()) & (data["HOUSE_NUMBER"].isna()), "HOUSE_NUMBER"] =
↳data["ADDRESS"].str.split(" ").str.get(0).str.split().str.get(0)
```

```
[ ]: data.loc[data["ADDRESS"] == "98A-98G Discala Ln", "HOUSE_NUMBER"] = "98A-98G"
```

Zamiana skrótów kierunków na pełne nazwy, usunięcie zbędnych znaków w nazwach ulic, zamiana liter na małe, usunięcie skrótów i pojedynczych liter w nazwach ulic. Usunięcie części kolumn i duplikatów.

```
[ ]: import re
#Zamiana skrótów na pełne nazwy
```

```

data.loc[data["STREET"].str.contains(r'\bE\b'), "STREET"] = data["STREET"].str.
    ↪replace("e", "east")
data.loc[data["STREET"].str.contains(r'\bW\b'), "STREET"] = data["STREET"].str.
    ↪replace("w", "west")
data.loc[data["STREET"].str.contains(r'\bS\b'), "STREET"] = data["STREET"].str.
    ↪replace("s", "south")
data.loc[data["STREET"].str.contains(r'\bN\b'), "STREET"] = data["STREET"].str.
    ↪replace("n", "north")

#Czyszczenie nazw ulic
data.loc[data["STREET"].str.contains(r' #.*'), "STREET"] = data["STREET"].str.
    ↪replace(r' #.*', "")
data["STREET"] = data["STREET"].str.replace(".", "")
data["STREET"] = data["STREET"].str.lower()
data["STREET"].unique()

def clean_street_names(street_name):

    parts = street_name.split(' ')
    cleaned_parts = [part for part in parts if re.match(r'^\d*(?:th|st|rd|nd)?
    ↪$', part) or not any(c.isdigit() for c in part) ]
    cleaned_parts = [part for part in cleaned_parts if not part.isdigit() and
    ↪len(part) != 1]

    last_part = cleaned_parts[-1] if cleaned_parts else None
    if last_part and any(character.isdigit() for character in last_part):
        cleaned_parts.remove(last_part)

    return ' '.join(cleaned_parts)

data['STREET'] = data['STREET'].apply(clean_street_names)

data['STREET']

```

```

[ ]: 0      east 55th street
      2      sinclair avenue
      3      55th street
      5      park place
      6      126th street
      ...
4796    80th street
4797    62nd drive
4798    21st street
4799    corona avenue
4800    neptune avenue
Name: STREET, Length: 4507, dtype: object

```

```
[ ]: data.columns
```

```
[ ]: Index(['BROKERTITLE', 'TYPE', 'PRICE', 'BEDS', 'BATH', 'PROPERTYSQFT',  
          'ADDRESS', 'STATE', 'ADMINISTRATIVE_AREA_LEVEL_2', 'LOCALITY',  
          'SUBLOCALITY', 'STREET_NAME', 'LONG_NAME', 'FORMATTED_ADDRESS',  
          'LATITUDE', 'LONGITUDE', 'BROKER', 'ANNOUNCEMENT_TYPE',  
          'PROPERTYSQFT1000', 'BOROUGH', 'NEIGHBOURHOOD', 'STREET', 'POSTCODE',  
          'HOUSE_NUMBER'],  
          dtype='object')
```

```
[ ]: data.drop(columns = ['BROKERTITLE',  
                          'ADDRESS', 'STATE', 'ADMINISTRATIVE_AREA_LEVEL_2', 'LOCALITY',  
                          'SUBLOCALITY', 'STREET_NAME', 'LONG_NAME', 'FORMATTED_ADDRESS',  
                          'PROPERTYSQFT1000'], inplace=True)
```

```
[ ]: category_columns = list(set(data.columns) - set(numeric_columns))
```

```
[ ]: for column in category_columns:  
      data[column] = data[column].str.lower()
```

```
[ ]: data.columns
```

```
[ ]: Index(['TYPE', 'PRICE', 'BEDS', 'BATH', 'PROPERTYSQFT', 'LATITUDE',  
          'LONGITUDE', 'BROKER', 'ANNOUNCEMENT_TYPE', 'BOROUGH', 'NEIGHBOURHOOD',  
          'STREET', 'POSTCODE', 'HOUSE_NUMBER'],  
          dtype='object')
```

```
[ ]: print('Duplicated rows: ', data.duplicated().sum())  
      data.drop_duplicates(inplace=True)
```

Duplicated rows: 1

```
[ ]: data.to_excel('data.xlsx')
```

GEOPY Pobieranie danych adresowych z geolokatora na podstawie współrzędnych geograficznych w celu uzupełnienia danych w kolumnie neighbourhood.

```
[ ]: #kod został zakomentowany aby dane nie zostały utracone przy uruchamianiu kodu  
      """from geopy.geocoders import Nominatim  
      import pandas as pd  
  
      data_api = pd.read_excel('data.xlsx')  
  
      data_api['API_ADDRESS_NAME'] = None  
      data_api['API_ADDRESS'] = None  
      api_keys = set()
```

```

num = 0
geolocator = Nominatim(user_agent="key_for_library1")

api_keys = set()
try:
    # Pętla służąca do pobrania danych dla każdego wiersza w ramce.
    for ind, row in data_api.iterrows():
        latitude = row['LATITUDE']
        longitude = row['LONGITUDE']
        address = geolocator.reverse(f"{latitude},{longitude}")
        data_api.at[ind, 'API_ADDRESS_NAME'] = address.raw['display_name']
        data_api.at[ind, 'API_ADDRESS'] = address.raw
        keys = set(address.raw['address'].keys())
        api_keys.update(keys)
        print(ind)
except Exception as e:
    data_api.to_excel('api.xlsx')

data_api.to_excel('api.xlsx')

data_api = pd.read_excel('api.xlsx')

geolocator = Nominatim(user_agent="key_for_library2")

try:
    # Pętla służąca do pobrania danych dla każdego wiersza w ramce.
    for ind, row in data_api.iterrows():
        if ind >= 2010:
            latitude = row['LATITUDE']
            longitude = row['LONGITUDE']
            address = geolocator.reverse(f"{latitude},{longitude}")
            data_api.at[ind, 'API_ADDRESS_NAME'] = address.raw['display_name']
            data_api.at[ind, 'API_ADDRESS'] = address.raw
            keys = set(address.raw['address'].keys())
            api_keys.update(keys)
            print(ind)
except Exception as e:
    data_api.to_excel('api_2.xlsx')

data_api.to_excel('api_2.xlsx')
print('Ilość kluczy: ', len(api_keys))
print('Klucze: ', api_keys)
"""

```

```

[ ]: 'from geopy.geocoders import Nominatim\nimport pandas as pd\n\ndata_api =
pd.read_excel(\'data.xlsx\')\n\ndata_api[\'API_ADDRESS_NAME\'] =

```

```

None\data_api['API_ADDRESS'] = None\napi_keys = set()\n\nnum = 0\ngeolocator
= Nominatim(user_agent="key_for_library1")\n\napi_keys = set()\ntry:\n    #
Pętla służąca do pobrania danych dla każdego wiersza w ramce.\n    for ind, row
in data_api.iterrows():\n        latitude = row['LATITUDE']\n        longitude
= row['LONGITUDE']\n        address =
geolocator.reverse(f"{latitude},{longitude}")\n        data_api.at[ind,
'API_ADDRESS_NAME'] = address.raw['display_name']\n        data_api.at[ind,
'API_ADDRESS'] = address.raw\n        keys =
set(address.raw['address'].keys())\n        api_keys.update(keys)\n
print(ind)\nexcept Exception as e:\n
data_api.to_excel('api.xlsx')\n\nndata_api.to_excel('api.xlsx')\n\nndata_api
= pd.read_excel('api.xlsx')\n\ngeolocator =
Nominatim(user_agent="key_for_library2")\n\ntry:\n    # Pętla służąca do
pobrania danych dla każdego wiersza w ramce.\n    for ind, row in
data_api.iterrows():\n        if ind >= 2010:\n            latitude =
row['LATITUDE']\n            longitude = row['LONGITUDE']\n
address = geolocator.reverse(f"{latitude},{longitude}")\n
data_api.at[ind, 'API_ADDRESS_NAME'] = address.raw['display_name']\n
data_api.at[ind, 'API_ADDRESS'] = address.raw\n            keys =
set(address.raw['address'].keys())\n            api_keys.update(keys)\n
print(ind)\nexcept Exception as e:\n    data_api.to_excel('api_2.xlsx')\n\ndat
a_api.to_excel('api_2.xlsx')\n\nprint('Ilość kluczy:
',len(api_keys))\n\nprint('Klucze: ',api_keys)\n'

```

```

[ ]: import pandas as pd
import json

api_keys = ['neighbourhood', 'postcode', 'road', 'house_number']

data = pd.read_excel('api_2.xlsx')


#strukturyzowanie słownika danych uzyskanych dzięki API
for key in api_keys:
    for ind in data.index:
        try:
            row = data.at[ind, 'API_ADDRESS']
            row = str(row).replace("'", '"')
            row = json.loads(row)
            keys = list(row['address'].keys())
            if key in keys:
                data.at[ind, key] = row['address'][key].lower()
            else:
                data.at[ind, key] = ''
        except Exception as err:
            continue
data['NEIGHBOURHOOD'] = data['NEIGHBOURHOOD'].fillna(data['neighbourhood'])
data['NEIGHBOURHOOD'] = data['NEIGHBOURHOOD'].str.lower()

```

```
data = data[['TYPE', 'PRICE', 'BEDS', 'BATH', 'PROPERTYSQFT', 'LATITUDE', 'LONGITUDE', 'BROKER', 'ANNOUNCEMENT_TYPE', 'BOROUGH', 'STREET', 'POSTCODE', 'HOUSE_NUMBER', 'NEIGHBOURHOOD']]
```

```
#data.to_excel('cat_api.xlsx')
```

```
[ ]: plt.figure(figsize=(10, 8))
sns.scatterplot(data=data, x='LONGITUDE', y='LATITUDE', hue='NEIGHBOURHOOD', size='PROPERTYSQFT', sizes=(20, 200), legend=False)
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.show()
```



```
[ ]: data['NEIGHBOURHOOD'] = data['NEIGHBOURHOOD'].fillna('')
data['NEIGHBOURHOOD'].value_counts()
```

```
[ ]: NEIGHBOURHOOD
prospect heights      669
flushing              202
```

```

manhattan community board 8    180
jamaica                        99
forest hills                   88
...
new brighton                  1
university heights            1
kips bay                      1
belle harbor                  1
chinatown                     1
Name: count, Length: 282, dtype: int64

```

```
[ ]: data['NEIGHBOURHOOD'] = data['NEIGHBOURHOOD'].apply(lambda x: x if all(borough_
↳not in x for borough in listOfBoroughs) and 'bronx' not in x else '')
```

```
[ ]: data['NEIGHBOURHOOD'] = data['NEIGHBOURHOOD'].replace('', None)
data['NEIGHBOURHOOD'].value_counts()
```

```
[ ]: NEIGHBOURHOOD
prospect heights    669
flushing           202
jamaica             99
forest hills       88
streetaten island   73
...
westchester square    1
vinegar hill          1
parkville             1
peck slip             1
chinatown             1
Name: count, Length: 264, dtype: int64

```

```
[ ]: data.isna().sum()
```

```
[ ]: TYPE           0
PRICE              0
BEDS               0
BATH               0
PROPERTYYSQFT      0
LATITUDE           0
LONGITUDE          0
BROKER             0
ANNOUNCEMENT_TYPE  0
BOROUGH            0
STREET             0
POSTCODE           0
HOUSE_NUMBER       3
NEIGHBOURHOOD      791

```


dtype: int64

```
[ ]: category_columns = list(set(data.columns) - numeric_columns)
numeric_columns = list(numeric_columns)
data[list(category_columns)].nunique()
```

```
[ ]: BROKER          947
      TYPE           12
      POSTCODE      178
      HOUSE_NUMBER  2397
      ANNOUNCEMENT_TYPE  5
      STREET       1714
      NEIGHBOURHOOD  264
      BOROUGH        5
      dtype: int64
```

```
[ ]: sns.heatmap(data[numeric_columns].corr(), annot=True, cmap='Spectral',
                 linewidths=0.5,fmt=".2f", vmax=1, vmin=-1)
```

```
[ ]: <Axes: >
```




```
[ ]: #data.to_excel('ssdata.xlsx')
```

```
[ ]: # Dopasowanie i przeskalowanie danych dla danej kolumny
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
numeric_columns.remove('PRICE')

mm_scalers = {}
for column in numeric_columns:
    mm_scalers[column] = MinMaxScaler()
    data[column] = mm_scalers[column].fit_transform(data[column].values.
↪reshape(-1, 1))
```


5

```
[ ]: #Zamiana kategorii na wartości liczbowe
from sklearn.preprocessing import LabelEncoder
label_encoders = {}
for column in category_columns:
    label_encoders[column] = LabelEncoder()
    data[column] = label_encoders[column].fit_transform(data[column])
```

```
[ ]: for column in category_columns:
    for class_index, class_name in enumerate(label_encoders[column].classes_):
        print(f"{class_name}: {class_index}")
```

```
[ ]: plt.figure(figsize=(12, 10))
sns.heatmap(data.corr(), annot=True, cmap='Spectral', linewidths=0.5,fmt=".2f",
↪vmax=1, vmin=-1)
```

```
[ ]: <Axes: >
```



Utworzenie macierzy korelacji dla wszystkich kolumn.

Stworzenie data_relevant, zawierającej kolumny z dataframe data, których bezwzględna wartość korelacji z cenami nieruchomości wynosi co najmniej 0.03.

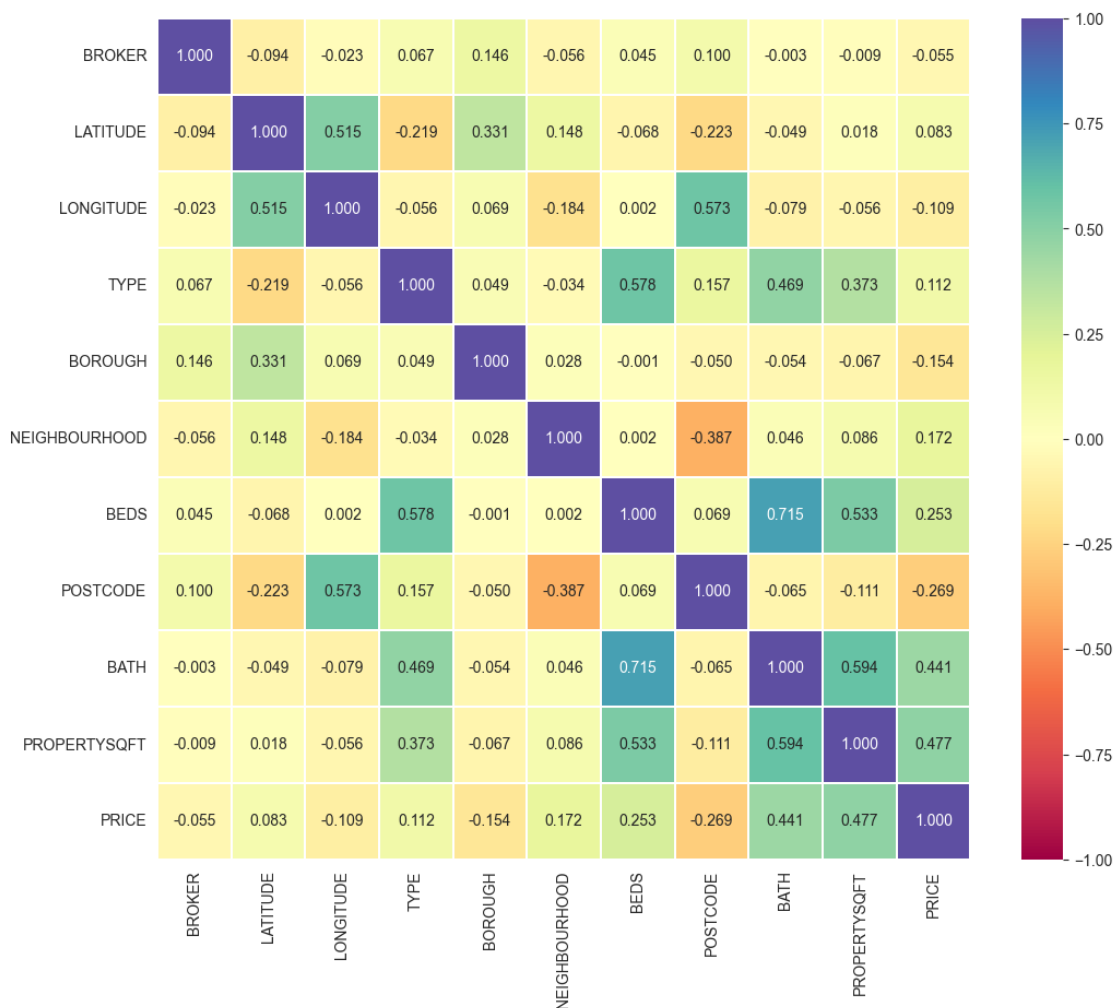
```
[ ]: corr_matrix = data.corr()
# Lista przechowująca pary kolumn, których współczynnik korelacji przekracza 0.8
pairs = []
for i in range(len(corr_matrix.columns)):
    for j in range(i+1, len(corr_matrix.columns)):
        if abs(corr_matrix.iloc[i, j]) > 0.8:
            pairs.append((corr_matrix.columns[i], corr_matrix.columns[j], corr_matrix.
↪ iloc[i, j]))

for pair in pairs:
    print(pair)
```

```
[ ]: #Filtrowanie DataFrame tak aby zawierał tylko kolumny silnie skorelowane z ceną
corr_matrix = data.corr()
corr_df = pd.DataFrame(abs(corr_matrix['PRICE'].drop('PRICE')).
    ↪sort_values(ascending=True))
relevant_columns = corr_df[abs(corr_df['PRICE'])>=0.05].index
relevant_columns = list(relevant_columns)
relevant_columns.append('PRICE')
data_relevant = data[relevant_columns]
```

```
[ ]: plt.figure(figsize=(12, 10))
sns.heatmap(data_relevant.corr(), annot=True, cmap='Spectral', linewidths=0.
    ↪3,fmt=".3f", vmax=1, vmin=-1)
```

```
[ ]: <Axes: >
```



```
[ ]: #data.to_excel('clean_data.xlsx')
data_relevant.to_excel('clean_data_relevant.xlsx')
```

Analiza głównych składowych PCA na danych: określenie liczby składowych głównych, stworzenie listy składowych głównych i nazw kolumn dla wynikowych składowych głównych, stworzenie instancji klasy PCA. Stworzenie ramki danych zawierającą przekształcone wartości zmiennych za pomocą PCA.

```
[ ]: from sklearn.decomposition import PCA
X = data_relevant.drop(columns=['PRICE'])
y = data_relevant['PRICE']

n_comp = 3
col_names = ['feature_'+str(i) for i in range(0,n_comp)]
pca = PCA(n_components=n_comp)
data_relevant_pca = pd.DataFrame(pca.fit_transform(X), columns = col_names)
print(pca.explained_variance_ratio_)
data_relevant_pca
```

```
[0.88876572 0.086828 0.02420772]
```

```
[ ]:      feature_0  feature_1  feature_2
0    -161.996266 -115.604465  26.933015
1     405.600518 -113.137165   1.529599
2    -221.973645 -114.140947  26.179353
3     408.315389 -26.271092 -46.062553
4    -161.897677 -113.999705  22.199013
...
4501 -221.522691 -106.138266   5.696170
4502  155.724427 -17.484736 -69.383974
4503 -162.175676 -118.498272  35.479340
4504 -145.177221  89.001470 -38.528559
4505  219.603802  97.280650   5.694335
```

```
[4506 rows x 3 columns]
```

```
[ ]: data_relevant_pca.to_excel('clean_data_relevant_pca.xlsx')
```

Planujemy użyć zarówno `clean_data_relevant_pca` oraz `clean_data_relevant` do trenowania modelu w następnej części projektu.

model-podstawowy-benchmark

June 9, 2024

```
[1]: import pandas as pd
cleaned_data = pd.read_excel('clean_data_relevant.xlsx')
cleaned_data = cleaned_data[['BROKER', 'LATITUDE', 'LONGITUDE', 'TYPE',
↪ 'BOROUGH', 'NEIGHBOURHOOD', 'BEDS', 'POSTCODE', 'BATH', 'PROPERTYSQFT',
↪ 'PRICE']]
```

```
[2]: cleaned_data
```

```
[2]:
```

	BROKER	LATITUDE	LONGITUDE	TYPE	BOROUGH	NEIGHBOURHOOD	BEDS	\
0	277	0.633396	0.505918	2	1	264	0.166667	
1	844	0.102276	0.103390	6	3	259	0.333333	
2	217	0.633742	0.505683	2	1	264	0.250000	
3	844	0.423098	0.534539	6	0	192	0.416667	
4	277	0.750035	0.556240	2	1	264	0.166667	
...	
4501	217	0.665089	0.539708	0	1	263	0.083333	
4502	591	0.563894	0.713574	0	2	196	0.083333	
4503	277	0.596189	0.453402	0	1	264	0.083333	
4504	288	0.588660	0.690689	2	2	91	0.166667	
4505	653	0.192653	0.512337	0	0	62	0.083333	
	POSTCODE	BATH	PROPERTYSQFT	PRICE				
0	18	0.2	0.134948	315000				
1	54	0.2	0.205882	260000				
2	18	0.1	0.024798	69000				
3	127	0.2	0.435294	690000				
4	23	0.2	0.225399	899500				
...				
4501	40	0.1	0.225399	599000				
4502	147	0.1	0.225399	245000				
4503	9	0.1	0.225399	1275000				
4504	146	0.1	0.049020	598125				
4505	114	0.1	0.059977	349000				

```
[4506 rows x 11 columns]
```

```
[3]: from sklearn.model_selection import train_test_split

X = cleaned_data.drop(columns='PRICE')
y = cleaned_data['PRICE']

[4]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
↳random_state=42)

[5]: from sklearn.linear_model import LinearRegression, HuberRegressor

reg = LinearRegression().fit(X_train, y_train)

print('Współczynniki: ', reg.coef_)
print('Wyraz wolny', reg.intercept_)

y_pred = reg.predict(X_test)

from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error,
↳median_absolute_error
import numpy as np

print("\nR-squared:", r2_score(y_test, y_pred))
print("Mean Absolute Error:", mean_absolute_error(y_test, y_pred))
print("Median Absolute Error:", median_absolute_error(y_test, y_pred))
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("Root Mean Squared Error:", np.sqrt(mean_squared_error(y_test, y_pred)))
```

```
Współczynniki: [ 4.82637713e+01  1.04212168e+06  4.80951990e+05 -5.98126037e+04
 -3.87799376e+05  2.53902786e+03 -2.46696958e+06 -1.27222931e+04
  8.99609911e+06  8.26434460e+06]
Wyraz wolny -625624.917808031
```

```
R-squared: 0.35106361164412225
Mean Absolute Error: 1336828.6355151676
Median Absolute Error: 774706.2717598878
Mean Squared Error: 6797303881835.929
Root Mean Squared Error: 2607163.953769676
```


Wartość współczynnika determinacji (R^2) wynosi około 0.35, co oznacza, że około 35% zmienności zmiennej PRICE jest wyjaśniona przez nasz model regresji liniowej. Idealna wartość tej miary wynosi 1, więc nasz model sprawuje się umiarkowanie w przewidywaniu nowych wartości.

Wysokie wartości średniego absolutnego błędu, mediany absolutnego błędu oraz błędu średniokwadratowego sugerują, że model ma ograniczoną jakość predykcyjną.

RMSE, czyli pierwiastek średniokwadratowy błędu, wynoszący około 2.6 miliona, oznacza że wartości przewidywane przez nasz model średnio odbiegają od rzeczywistych wartości o 2.6 miliona.

```
[6]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, color='blue', label='Wartości Rzeczywiste a
↳Przewidywane')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], '--',
↳color='red', label='Ideal')
plt.title('Wartości Rzeczywiste a Przewidywane')
plt.xlabel('Rzeczywiste')
plt.ylabel('Przewidywane')
plt.legend()
plt.show()
```



Z wykresu możemy wnioskować, że model średnio sobie radzi z przewidywaniem wyższych wartości. Może być to spowodowane niewystarczającą liczbą przypadków ogłoszeń z cenami >1.5 miliona lub wybraniem złego modelu.

```
[7]: from sklearn.model_selection import cross_val_score, RepeatedKFold
from sklearn.pipeline import Pipeline
cv = RepeatedKFold(n_splits=10, n_repeats=10, random_state=1)
pipe_linear = Pipeline([
    ('linear', LinearRegression(fit_intercept=True))
])
```

```

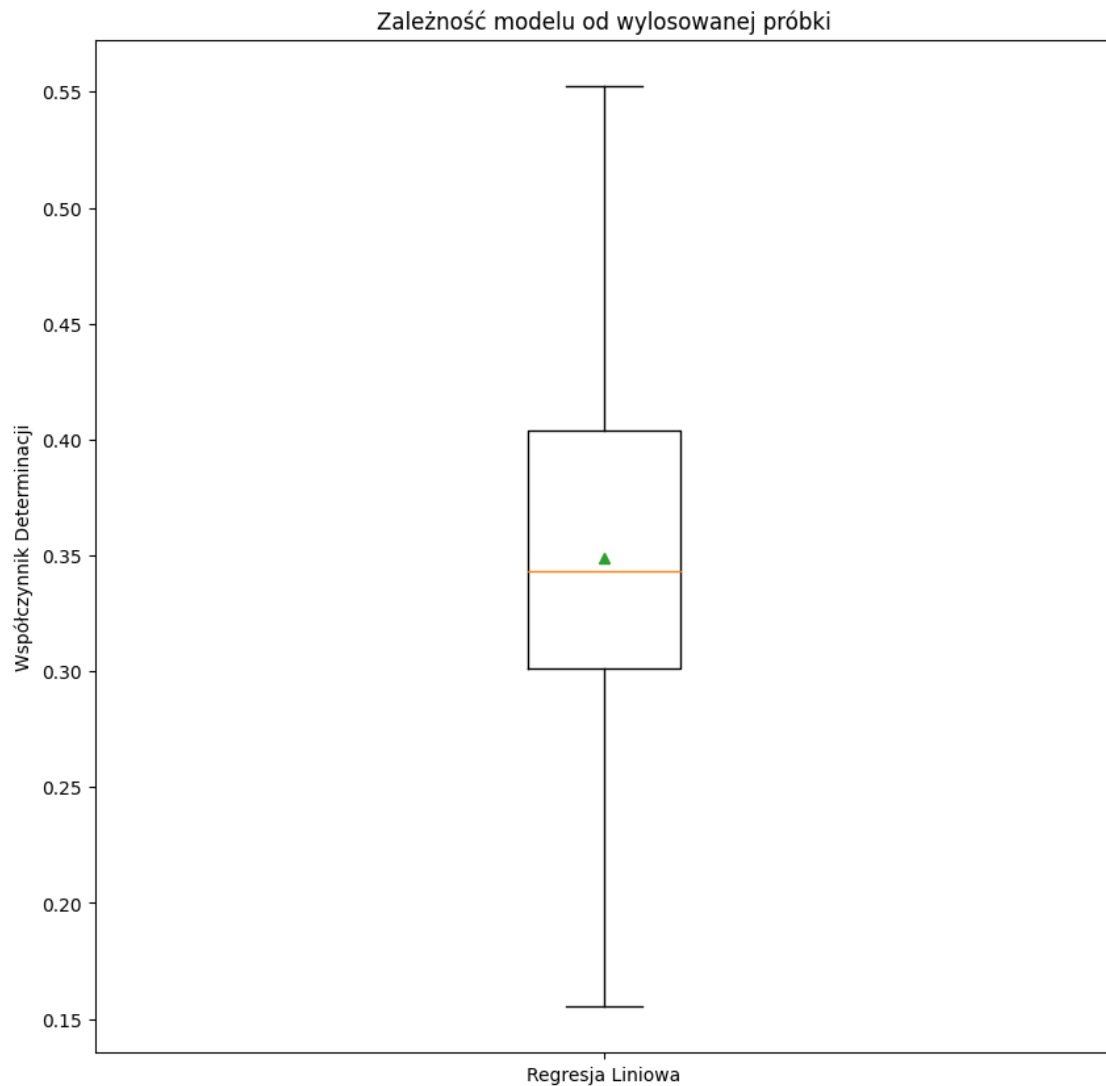
model_linear_scores = cross_val_score(pipe_linear, X, y, scoring='r2', cv = cv)
print("Cross Validation Score: ", model_linear_scores)
plt.figure(figsize=(10, 10))
plt.boxplot([model_linear_scores], labels=['Regresja Liniowa'], showmeans=True)
plt.title('Zależność modelu od wylosowanej próbki')
plt.ylabel('Współczynnik Determinacji')
plt.show()

```

```

Cross Validation Score:  [0.3048472  0.42354097 0.34307738 0.32622976 0.40298448
0.33415731
0.35389578 0.43982456 0.34043789 0.23120915 0.48988149 0.24331674
0.49782386 0.28851387 0.38692206 0.31823726 0.5076094  0.24137524
0.20087974 0.3717199  0.30186297 0.36090288 0.36091703 0.47216805
0.32848872 0.40213783 0.28580038 0.3227488  0.31614519 0.31932509
0.27444135 0.34849079 0.42773341 0.55234023 0.23181206 0.34301315
0.18369045 0.37432225 0.30588735 0.45629086 0.32173945 0.35631834
0.33146502 0.30836542 0.48871109 0.33438085 0.32668307 0.21656633
0.40950851 0.46846566 0.24628029 0.35716169 0.38562963 0.36004001
0.31114419 0.44091716 0.52581639 0.24018291 0.21947161 0.36402104
0.33655075 0.28235667 0.48104963 0.38214224 0.44749608 0.2721098
0.3114152  0.38592847 0.32025077 0.27704416 0.29393787 0.37346854
0.37172132 0.40549431 0.41534919 0.32788038 0.44813303 0.26441936
0.20789418 0.31583986 0.18180707 0.37524692 0.28147825 0.4478128
0.37382166 0.35915256 0.42621198 0.28800913 0.38681281 0.37627307
0.40969144 0.29843979 0.4057372  0.46327359 0.33209954 0.30475667
0.48181861 0.15528443 0.2295418  0.35885797]

```

Na podstawie wyników walidacji krzyżowej możemy powiedzieć, że model nie jest stabilny, ponieważ wyniki walidacji krzyżowej różnią się dosyć znacząco dla różnych podziałów danych. Przykładowo najwyższy wynik R^2 to 0.55 a najniższy około 0.15.

```
[8]: from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LinearRegression

pipe_linear = Pipeline([
    ('linear', LinearRegression(fit_intercept=True))
])

param_grid = {
    'linear_fit_intercept': [True, False],
    'linear_n_jobs': [None, 1, 2, 4, 8],
```

```

        'linear__positive': [False, True]
    }

cv = RepeatedKFold(n_splits=10, n_repeats=10, random_state=1)

# Perform grid search
grid_search = GridSearchCV(pipe_linear, param_grid, cv=cv, scoring='r2')
grid_search.fit(X, y)

# Get the best parameters and best score
best_params = grid_search.best_params_
best_score = grid_search.best_score_

print("Best Parameters:", best_params)
print("Best Score:", best_score)

```

```

Best Parameters: {'linear__fit_intercept': True, 'linear__n_jobs': None,
'linear__positive': False}
Best Score: 0.3488448065360991

```

```

[9]: from statsmodels.stats.outliers_influence import variance_inflation_factor
vif_x=X.copy()
vif_data = pd.DataFrame()
vif_data["feature"] = vif_x.columns

vif_data["VIF"] = [variance_inflation_factor(vif_x.values, i)
                    for i in range(len(vif_x.columns))]

print(vif_data)

```

	feature	VIF
0	BROKER	3.424444
1	LATITUDE	24.088323
2	LONGITUDE	46.314859
3	TYPE	3.725739
4	BOROUGH	3.215143
5	NEIGHBOURHOOD	4.328808
6	BEDS	9.242304
7	POSTCODE	13.026365
8	BATH	8.219750
9	PROPERTYSQFT	5.652389

Wysokie wartości VIF dla zmiennych BEDS, POSTCODE, BATH, PROPERTYSQFT, LATITUDE oraz LONGITUDE wskazują na ich silną korelację z innymi zmiennymi w zbiorze danych. Z uwagi na to, że wysokie wartości VIF są zazwyczaj niepożądane i mogą świadczyć o możliwej wieloliniowości modelu, rozważenie działań mających na celu zmniejszenie korelacji, np. przez zastosowanie analizy głównych składowych (PCA), może być wskazane.

```
[10]: from sklearn.decomposition import PCA

n_comp = 6
col_names = ['feature_'+str(i) for i in range(0,n_comp)]
pca = PCA(n_components=n_comp)
data_relevant_pca = pd.DataFrame(pca.fit_transform(X), columns = col_names)

data_relevant_pca
```

```
[10]:      feature_0  feature_1  feature_2  feature_3  feature_4  feature_5
0   -161.996266 -115.604465  26.933015  -1.502278  -0.524361  -0.060368
1    405.600518 -113.137165   1.529599   1.664794   1.029621   0.612129
2   -221.973645 -114.140947  26.179353  -1.457868  -0.480821  -0.040537
3    408.315389  -26.271092 -46.062553   0.790193  -1.800866  -0.051927
4   -161.897677 -113.999705  22.199013  -1.565854  -0.508710  -0.188432
...
4501 -221.522691 -106.138266   5.696170  -3.734768  -0.407511  -0.058016
4502  155.724427  -17.484736 -69.383974  -5.250187   0.523221   0.006469
4503 -162.175676 -118.498272  35.479340  -3.388526  -0.508557   0.015195
4504 -145.177221   89.001470 -38.528559  -2.844816   0.720788   0.006368
4505  219.603802   97.280650   5.694335  -4.703789  -1.571506   0.261213

[4506 rows x 6 columns]
```

```
[11]: X_train, X_test, y_train, y_test = train_test_split(data_relevant_pca, y,
↳ test_size=0.3, random_state=42)
```

```
[12]: from statsmodels.stats.outliers_influence import variance_inflation_factor
vif_x=data_relevant_pca.copy()

vif_data = pd.DataFrame()
vif_data["feature"] = vif_x.columns

vif_data["VIF"] = [variance_inflation_factor(vif_x.values, i)
                    for i in range(len(vif_x.columns))]

print(vif_data)
```

```
      feature  VIF
0  feature_0  1.0
1  feature_1  1.0
2  feature_2  1.0
3  feature_3  1.0
4  feature_4  1.0
5  feature_5  1.0
```

Wartości VIF wynoszące 1.0 dla wszystkich zmiennych wskazują na brak korelacji między nimi. Taka sytuacja jest korzystna dla stabilności wybranego modelu.

```
[13]: reg_pca = LinearRegression().fit(X_train, y_train)

print('Współczynniki modelu: ', reg_pca.coef_)
print('Wyraz wolny: ', reg_pca.intercept_)

y_pred = reg_pca.predict(X_test)

print("\nR-squared:", r2_score(y_test, y_pred))
print("Mean Absolute Error:", mean_absolute_error(y_test, y_pred))
print("Median Absolute Error:", median_absolute_error(y_test, y_pred))
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("Root Mean Squared Error:", np.sqrt(mean_squared_error(y_test, y_pred)))
```

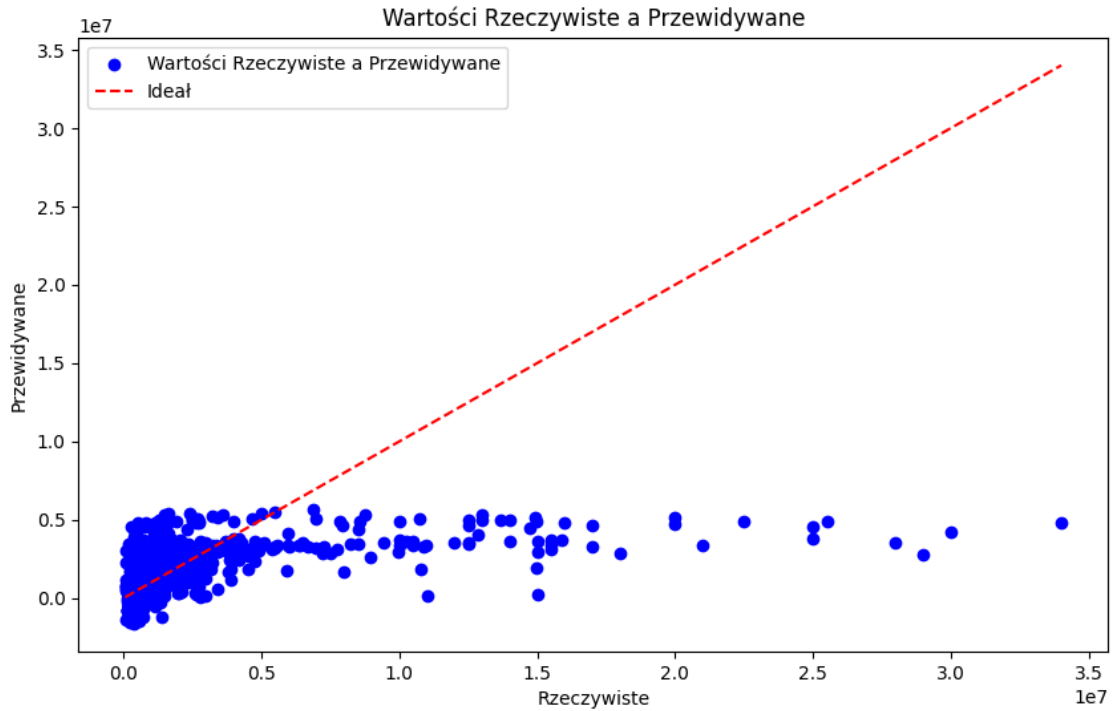
```
Współczynniki modelu:  [-8.00408231e+02 -8.48885438e+03  1.45686807e+04
 1.32674542e+05
 -4.64438459e+05 -2.75789608e+06]
Wyraz wolny:  1710548.1834018864
```

```
R-squared: 0.1867849813104121
Mean Absolute Error: 1472007.5710042594
Median Absolute Error: 833256.5642988225
Mean Squared Error: 8518045377776.891
Root Mean Squared Error: 2918569.06338995
```

[13]:

Model wytrenowany na danych po transformacji PCA wydaje się mieć ograniczoną zdolność do wyjaśniania zmienności cechy PRICE na podstawie dostarczonego zbioru - wynik współczynnika determinacji wynosi około 0.15. Ponadto, wysokie wartości błędów wskazują na znaczną rozbieżność między przewidywanymi a rzeczywistymi wartościami zmiennej PRICE.


```
[14]: plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, color='blue', label='Wartości Rzeczywiste a_
↳Przewidywane')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], '--',
↳color='red', label='Ideal')
plt.title('Wartości Rzeczywiste a Przewidywane')
plt.xlabel('Rzeczywiste')
plt.ylabel('Przewidywane')
plt.legend()
plt.show()
```



```
[15]: cv = RepeatedKfold(n_splits=5, n_repeats=5, random_state=1)
pipe_linear = Pipeline([
    ('linear', LinearRegression(fit_intercept=True))
])

model_linear_scores = cross_val_score(pipe_linear, data_relevant_pca, y,
    scoring='r2', cv = cv)
print("Cross Validation Score: ", model_linear_scores)
plt.figure(figsize=(10, 10))
plt.boxplot([model_linear_scores], labels=['Linear Regression'], showmeans=True)
plt.title('Jakość modelu')
plt.ylabel('Negative Mean Absolute Error')
plt.show()
```

```
Cross Validation Score: [0.22456969 0.12135091 0.16043687 0.2207802  0.1433581
0.17514762
0.19745571 0.16164154 0.15713593 0.13710569 0.16294432 0.19592756
0.21469362 0.13717919 0.12834579 0.14219564 0.20406104 0.14507362
0.13142583 0.17862491 0.18951502 0.14731603 0.15102189 0.13690514
0.21530505]
```



dwa-dodatkowe-modele

June 9, 2024

1 Regresja

```
[20]: from matplotlib import pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
from sklearn.model_selection import cross_val_score
from sklearn import metrics
cleaned_data = pd.read_excel('clean_data_relevant.xlsx')
cleaned_data = cleaned_data[['BROKER', 'LATITUDE', 'LONGITUDE', 'TYPE',
    ↳ 'BOROUGH', 'NEIGHBOURHOOD', 'BEDS', 'POSTCODE', 'BATH', 'PROPERTYSQFT',
    ↳ 'PRICE']]
```

```
[21]: from sklearn.model_selection import train_test_split

X = cleaned_data.drop(columns='PRICE')
y = cleaned_data['PRICE']
#y = np.round(y, decimals=-2) #przy cenach posiadłości część dziesiętna nie ma
    ↳ aż tak dużo znaczenia. Zakładając że ostatnie 2 miejsca wynoszą 99
    ↳ to nawet dla minimalnej wartości w zbiorze danych stanowi to jedynie 0.2%
```

```
[22]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    ↳ random_state=42)
trained_models = []
```

1.1 Modele

1.1.1 Linear Model - model podstawowy

```
[23]: from sklearn.linear_model import LinearRegression, HuberRegressor

reg = LinearRegression().fit(X_train, y_train)

print('Współczynniki: ', reg.coef_)
print('Wyraz wolny', reg.intercept_)

y_pred = reg.predict(X_test)
```

```

cv_score = cross_val_score(estimator = reg, X = X_train, y = y_train, cv = 20)

RMSE = np.sqrt(metrics.mean_squared_error(y_test, y_pred))
R2 = reg.score(X_test, y_test)

print('RMSE:', round(RMSE,4))
print('R2:', round(R2,4))
print("Cross Validated R2: ", list(map(lambda x: round(x,4) ,cv_score)))
print("Mean Cross Validated R2: ", round(cv_score.mean(),4) )
print("Min Cross Validated R2: ", round(cv_score.min(),4) )
print("Max Cross Validated R2: ", round(cv_score.max(),4) )
trained_models.append(['Linear Model', round(RMSE,2) , round(R2,4),
↳round(cv_score.mean(),4), round(cv_score.min(),4), round(cv_score.max(),4),
↳list(map(lambda x: round(x,4) ,cv_score)) ])

```

Współczynniki: [4.82637713e+01 1.04212168e+06 4.80951990e+05 -5.98126037e+04
-3.87799376e+05 2.53902786e+03 -2.46696958e+06 -1.27222931e+04
8.99609911e+06 8.26434460e+06]

Wyraz wolny -625624.917808031

RMSE: 2607163.9538

R2: 0.3511

Cross Validated R2: [0.3458, 0.1537, 0.5281, 0.4193, 0.3508, 0.2672, 0.0612,
0.1882, 0.548, 0.2921, 0.1532, 0.4424, 0.0416, 0.2398, 0.4881, 0.5676, 0.4707,
0.5462, -0.0848, 0.2286]

Mean Cross Validated R2: 0.3124


Min Cross Validated R2: -0.0848

Max Cross Validated R2: 0.5676

```

[24]: plt.figure(figsize = (10,5))
sns.regplot(x=y_test,y=y_pred)
plt.title('Linear Regression for Housing dataset', fontsize = 20)
plt.show()

```

```
[25]: from sklearn.linear_model import LinearRegression, HuberRegressor

reg = HuberRegressor().fit(X_train, y_train)

print('Współczynniki: ', reg.coef_)
print('Wyraz wolny', reg.intercept_)

y_pred = reg.predict(X_test)

cv_score = cross_val_score(estimator = reg, X = X_train, y = y_train, cv = 20)

RMSE = np.sqrt(metrics.mean_squared_error(y_test, y_pred))
R2 = reg.score(X_test, y_test)

print('RMSE:', round(RMSE,4))
print('R2:', round(R2,4))
print("Cross Validated R2: ", list(map(lambda x: round(x,4) ,cv_score)))
print("Mean Cross Validated R2: ", round(cv_score.mean(),4) )
print("Min Cross Validated R2: ", round(cv_score.min(),4) )
print("Max Cross Validated R2: ", round(cv_score.max(),4) )
trained_models.append(['Huber Model', round(RMSE,2) , round(R2,4),
↳round(cv_score.mean(),4), round(cv_score.min(),4), round(cv_score.max(),4),
↳list(map(lambda x: round(x,4) ,cv_score)) ])
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:
ConvergenceWarning: lbfgs failed to converge (status=1):
```

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

```
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
```

Współczynniki: [-4.78448114e+01 5.52961796e+05 3.35432760e+05 6.80072102e+04
-2.09702782e+05 -1.03320073e+02 3.74172633e+05 -5.37332761e+03
5.26283624e+05 4.50102613e+05]

Wyraz wolny 765420.4355215558

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:

ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

```
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:

ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

```
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:

ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

```
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:

ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

```
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:

ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

```
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:

ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:
ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

```
https://scikit-learn.org/stable/modules/preprocessing.html
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:


```
https://scikit-learn.org/stable/modules/preprocessing.html
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)
```

```
RMSE: 3110867.2428
R2: 0.0761
Cross Validated R2: [0.0882, 0.2082, 0.1447, 0.0718, 0.0313, 0.0541, 0.2153,
0.1253, 0.04, 0.0194, 0.0095, 0.1026, 0.1088, 0.0383, 0.0795, 0.1062, 0.0581,
0.1332, 0.2582, 0.0244]
Mean Cross Validated R2: 0.0959
Min Cross Validated R2: 0.0095
Max Cross Validated R2: 0.2582

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_huber.py:342:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
self.n_iter_ = _check_optimize_result("lbfgs", opt_res, self.max_iter)

```
[26]: plt.figure(figsize = (10,5))
sns.regplot(x=y_test,y=y_pred)
plt.title('Huber Regression for Housing dataset', fontsize = 20)
plt.show()
```



1.1.2 Forest Regressor

```
[27]: from sklearn.ensemble import RandomForestRegressor

forest_reg = RandomForestRegressor(n_estimators = 10, random_state = 0)
forest_reg.fit(X_train, y_train)
y_pred = forest_reg.predict(X_test)

cv_score = cross_val_score(estimator = forest_reg, X = X_train, y = y_train, cv=
    ↪= 20)
RMSE = np.sqrt(metrics.mean_squared_error(y_test, y_pred))
R2 = forest_reg.score(X_test, y_test)

print('RMSE:', round(RMSE,4))
print('R2:', round(R2,4))
print("Cross Validated R2: ", list(map(lambda x: round(x,4) ,cv_score)))
print("Mean Cross Validated R2: ", round(cv_score.mean(),4) )
print("Min Cross Validated R2: ", round(cv_score.min(),4) )
print("Max Cross Validated R2: ", round(cv_score.max(),4) )
trained_models.append(['Forest Regressor Model = 10 trees', round(RMSE,2) ,
    ↪round(R2,4), round(cv_score.mean(),4), round(cv_score.min(),4),
    ↪round(cv_score.max(),4), list(map(lambda x: round(x,4) ,cv_score)) ])
```

RMSE: 1940629.4766

R2: 0.6405

Cross Validated R2: [0.7517, -0.0809, 0.692, 0.5232, 0.7029, 0.1126, 0.1084,
0.2822, 0.8527, 0.4304, 0.2866, 0.2914, 0.2893, 0.6414, 0.6216, 0.7545, 0.7896,
0.5274, 0.0793, 0.4204]

Mean Cross Validated R2: 0.4538

Min Cross Validated R2: -0.0809

Max Cross Validated R2: 0.8527

```
[28]: from sklearn.ensemble import RandomForestRegressor

forest_reg_30 = RandomForestRegressor(n_estimators = 30, random_state = 0)
forest_reg_30.fit(X_train, y_train)
y_pred = forest_reg_30.predict(X_test)

cv_score = cross_val_score(estimator = forest_reg_30, X = X_train, y = y_train,
    ↪cv = 20)
RMSE = np.sqrt(metrics.mean_squared_error(y_test, y_pred))
R2 = forest_reg_30.score(X_test, y_test)

print('RMSE:', round(RMSE,4))
print('R2:', round(R2,4))
```

```

print("Cross Validated R2: ", list(map(lambda x: round(x,4) ,cv_score)))
print("Mean Cross Validated R2: ", round(cv_score.mean(),4) )
print("Min Cross Validated R2: ", round(cv_score.min(),4) )
print("Max Cross Validated R2: ", round(cv_score.max(),4) )
trained_models.append(['Forest Regressor Model = 30 trees', round(RMSE,2) ,
↳round(R2,4), round(cv_score.mean(),4), round(cv_score.min(),4),
↳round(cv_score.max(),4), list(map(lambda x: round(x,4) ,cv_score)) ])

```

RMSE: 1795072.1975

R2: 0.6924

Cross Validated R2: [0.7542, 0.1516, 0.7353, 0.445, 0.6769, 0.3038, 0.3995, 0.3781, 0.8728, 0.4241, 0.3019, 0.3444, 0.2986, 0.5409, 0.7355, 0.7796, 0.8565, 0.6355, 0.5446, 0.4774]

Mean Cross Validated R2: 0.5328


Min Cross Validated R2: 0.1516

Max Cross Validated R2: 0.8728

```

[29]: fig = plt.figure()
plt.scatter(y_test,y_pred)
fig.suptitle('y_test vs y_pred', fontsize=20)           # Plot heading
plt.xlabel('y_test', fontsize=18)                       # X-label
plt.ylabel('y_pred', fontsize=16)                       # Y-label
plt.show()

```



1.1.3 Decision Tree

```
[30]: from sklearn.tree import DecisionTreeRegressor
tree_reg = DecisionTreeRegressor(random_state=0)
tree_reg.fit(X_train, y_train)
y_pred = tree_reg.predict(X_test)

cv_score = cross_val_score(estimator = tree_reg, X = X_train, y = y_train, cv = 20)

RMSE = np.sqrt(metrics.mean_squared_error(y_test, y_pred))
R2 = tree_reg.score(X_test, y_test)

print('RMSE:', round(RMSE,4))
print('R2:', round(R2,4))
print("Cross Validated R2: ", list(map(lambda x: round(x,4) ,cv_score)))
print("Mean Cross Validated R2: ", round(cv_score.mean(),4) )
print("Min Cross Validated R2: ", round(cv_score.min(),4) )
```



```
print("Max Cross Validated R2: ", round(cv_score.max(),4) )
trained_models.append(['Decision Tree', round(RMSE,2) , round(R2,4),
↳round(cv_score.mean(),4), round(cv_score.min(),4), round(cv_score.max(),4),
↳list(map(lambda x: round(x,4) ,cv_score)) ])
```

RMSE: 3271182.2776

R2: -0.0216


Cross Validated R2: [0.5702, -2.84, -0.445, 0.4527, -0.24, 0.1602, 0.0184, 0.2744, 0.6666, -0.0912, 0.0288, 0.2898, -0.1282, 0.4911, 0.6577, 0.312, 0.2001, 0.2164, -0.8715, 0.2628]

Mean Cross Validated R2: -0.0007

Min Cross Validated R2: -2.84

Max Cross Validated R2: 0.6666

```
[31]: fig = plt.figure()
plt.scatter(y_test,y_pred)
fig.suptitle('y_test vs y_pred', fontsize=20)
plt.xlabel('y_test', fontsize=18)
plt.ylabel('y_pred', fontsize=16)
plt.show()
```



1.1.4 Ridge Regression

```
[32]: from sklearn.linear_model import Ridge

ridge_reg = Ridge(alpha=3, solver="cholesky")
ridge_reg.fit(X_train, y_train)
y_pred = ridge_reg.predict(X_test)

cv_score = cross_val_score(estimator = ridge_reg, X = X_train, y = y_train, cv=
    ↪ 20)
RMSE = np.sqrt(metrics.mean_squared_error(y_test, y_pred))
R2 = ridge_reg.score(X_test, y_test)

print('RMSE:', round(RMSE,4))
print('R2:', round(R2,4))
print("Cross Validated R2: ", list(map(lambda x: round(x,4) ,cv_score)))
print("Mean Cross Validated R2: ", round(cv_score.mean(),4) )
print("Min Cross Validated R2: ", round(cv_score.min(),4) )
print("Max Cross Validated R2: ", round(cv_score.max(),4) )
trained_models.append(['Ridge Model', round(RMSE,2) , round(R2,4),
    ↪round(cv_score.mean(),4), round(cv_score.min(),4), round(cv_score.max(),4),
    ↪list(map(lambda x: round(x,4) ,cv_score)) ])
```

RMSE: 2604068.0257

R2: 0.3526


Cross Validated R2: [0.3422, 0.1821, 0.5286, 0.4153, 0.3451, 0.2662, 0.0868,
0.1965, 0.5346, 0.2871, 0.1596, 0.4547, 0.0792, 0.2423, 0.4827, 0.5581, 0.4564,
0.5475, 0.0004, 0.2287]

Mean Cross Validated R2: 0.3197

Min Cross Validated R2: 0.0004

Max Cross Validated R2: 0.5581

```
[33]: fig = plt.figure()
plt.scatter(y_test,y_pred)
fig.suptitle('y_test vs y_pred', fontsize=20)
plt.xlabel('y_test', fontsize=18)
plt.ylabel('y_pred', fontsize=16)
plt.show()
```



1.1.5 XGBoost

```
[34]: from xgboost import XGBRegressor

XGBR = XGBRegressor(n_estimators=1000, max_depth=7, eta=0.1, subsample=0.8,
                    ↪ colsample_bytree=0.8)
XGBR.fit(X_train, y_train)
y_pred = XGBR.predict(X_test)

cv_score = cross_val_score(estimator = XGBR, X = X_train, y = y_train, cv = 20)
RMSE = np.sqrt(metrics.mean_squared_error(y_test, y_pred))
R2 = XGBR.score(X_test, y_test)

print('RMSE:', round(RMSE,4))
print('R2:', round(R2,4))
print("Cross Validated R2: ", list(map(lambda x: round(x,4) ,cv_score)))
print("Mean Cross Validated R2: ", round(cv_score.mean(),4) )
```

```

print("Min Cross Validated R2: ", round(cv_score.min(),4) )
print("Max Cross Validated R2: ", round(cv_score.max(),4) )
trained_models.append(['XGBRegressor', round(RMSE,2) , round(R2,4),
↳round(cv_score.mean(),4), round(cv_score.min(),4), round(cv_score.max(),4),
↳list(map(lambda x: round(x,4) ,cv_score)) ])

```

RMSE: 1879060.5311

R2: 0.6629

Cross Validated R2: [0.7585, -0.5486, 0.5267, 0.664, 0.6568, 0.4718, 0.4094, 0.414, 0.8921, 0.6182, 0.7159, -0.2433, 0.3732, 0.6535, 0.796, 0.7811, 0.8598, 0.7116, 0.5855, 0.3537]

Mean Cross Validated R2: 0.5225


Min Cross Validated R2: -0.5486

Max Cross Validated R2: 0.8921

```

[35]: fig = plt.figure()
plt.scatter(y_test,y_pred)
fig.suptitle('y_test vs y_pred', fontsize=20)
plt.xlabel('y_test', fontsize=18)
plt.ylabel('y_pred', fontsize=16)
plt.show()

```



1.1.6 SVR

```
[36]: from sklearn.svm import SVR
svr_clf = SVR(kernel = 'rbf')

svr_clf.fit(X_train, y_train)
y_pred = svr_clf.predict(X_test)

cv_score = cross_val_score(estimator = svr_clf, X = X_train, y = y_train, cv = 20)
RMSE = np.sqrt(metrics.mean_squared_error(y_test, y_pred))
R2 = svr_clf.score(X_test, y_test)

print('RMSE:', round(RMSE,4))
print('R2:', round(R2,4))
print("Cross Validated R2: ", list(map(lambda x: round(x,4) ,cv_score)))
print("Mean Cross Validated R2: ", round(cv_score.mean(),4) )
print("Min Cross Validated R2: ", round(cv_score.min(),4) )
```

```
print("Max Cross Validated R2: ", round(cv_score.max(),4) )
trained_models.append(['SVR', round(RMSE,2) , round(R2,4), round(cv_score.
↪mean(),4), round(cv_score.min(),4), round(cv_score.max(),4), list(map(lambda_
↪x: round(x,4) ,cv_score)) ])
```

RMSE: 3366019.8979

R2: -0.0817


Cross Validated R2: [-0.056, -0.1082, -0.1063, -0.0855, -0.0628, -0.0479,
-0.0372, -0.0459, -0.0884, -0.0641, -0.0652, -0.134, -0.1696, -0.0637, -0.0861,
-0.0799, -0.0733, -0.0934, -0.0609, -0.0494]

Mean Cross Validated R2: -0.0789

Min Cross Validated R2: -0.1696

Max Cross Validated R2: -0.0372

```
[37]: fig = plt.figure()
plt.scatter(y_test,y_pred)
fig.suptitle('SVR', fontsize=20)                    # Plot heading
plt.xlabel('y_test', fontsize=18)                  # X-label
plt.ylabel('y_pred', fontsize=16)                  # Y-label
plt.show()
```



1.1.7 Bayesian regression

```
[38]: from sklearn.linear_model import BayesianRidge
bayesian_reg = BayesianRidge()
bayesian_reg.fit(X_train, y_train)
y_pred = bayesian_reg.predict(X_test)

cv_score = cross_val_score(estimator = bayesian_reg, X = X_train, y = y_train,
    ↪cv = 20)
RMSE = np.sqrt(metrics.mean_squared_error(y_test, y_pred))
R2 = bayesian_reg.score(X_test, y_test)

print('RMSE:', round(RMSE,4))
print('R2:', round(R2,4))
print("Cross Validated R2: ", list(map(lambda x: round(x,4) ,cv_score)))
print("Mean Cross Validated R2: ", round(cv_score.mean(),4) )
print("Min Cross Validated R2: ", round(cv_score.min(),4) )
print("Max Cross Validated R2: ", round(cv_score.max(),4) )
trained_models.append(['Bayesian Reg', round(RMSE,2) , round(R2,4),
    ↪round(cv_score.mean(),4), round(cv_score.min(),4), round(cv_score.max(),4),
    ↪list(map(lambda x: round(x,4) ,cv_score)) ])
```

RMSE: 2606303.0194

R2: 0.3515


Cross Validated R2: [0.3452, 0.1593, 0.5284, 0.4186, 0.3498, 0.2671, 0.0668,
0.19, 0.5454, 0.2912, 0.1544, 0.4451, 0.0489, 0.2404, 0.4872, 0.5658, 0.4678,
0.5469, -0.0677, 0.2287]

Mean Cross Validated R2: 0.314

Min Cross Validated R2: -0.0677

Max Cross Validated R2: 0.5658

```
[39]: fig = plt.figure()
plt.scatter(y_test,y_pred)
fig.suptitle('Bayesian Ridge', fontsize=20)
plt.xlabel('y_test', fontsize=18)           # X-label
plt.ylabel('y_pred', fontsize=16)          # Y-label
plt.show()
```



1.2 Podsumowanie Regresji

```
[40]: trained_models = pd.DataFrame( trained_models, columns=['Model', 'RMSE', 'R2_
↪Score', 'Mean Cross Validated R2 Score', 'Min Cross Validated R2 Score', 'Max_
↪Cross Validated R2 Score', 'Cross Validated R2 Scores'])
trained_models
```

```
[40]:
```

	Model	RMSE	R2 Score \
0	Linear Model	2607163.95	0.3511
1	Huber Model	3110867.24	0.0761
2	Forest Regressor Model = 10 trees	1940629.48	0.6405
3	Forest Regressor Model = 30 trees	1795072.20	0.6924
4	Decision Tree	3271182.28	-0.0216
5	Ridge Model	2604068.03	0.3526
6	XGBRegressor	1879060.53	0.6629
7	SVR	3366019.90	-0.0817
8	Bayesian Reg	2606303.02	0.3515

	Mean Cross Validated R2 Score	Min Cross Validated R2 Score \
0	0.3124	-0.0848
1	0.0959	0.0095
2	0.4538	-0.0809
3	0.5328	0.1516
4	-0.0007	-2.8400
5	0.3197	0.0004
6	0.5225	-0.5486
7	-0.0789	-0.1696
8	0.3140	-0.0677

	Max Cross Validated R2 Score \
0	0.5676
1	0.2582
2	0.8527
3	0.8728
4	0.6666
5	0.5581
6	0.8921
7	-0.0372
8	0.5658


	Cross Validated R2 Scores
0	[0.3458, 0.1537, 0.5281, 0.4193, 0.3508, 0.267...
1	[0.0882, 0.2082, 0.1447, 0.0718, 0.0313, 0.054...
2	[0.7517, -0.0809, 0.692, 0.5232, 0.7029, 0.112...
3	[0.7542, 0.1516, 0.7353, 0.445, 0.6769, 0.3038...
4	[0.5702, -2.84, -0.445, 0.4527, -0.24, 0.1602,...
5	[0.3422, 0.1821, 0.5286, 0.4153, 0.3451, 0.266...
6	[0.7585, -0.5486, 0.5267, 0.664, 0.6568, 0.471...
7	[-0.056, -0.1082, -0.1063, -0.0855, -0.0628, -...
8	[0.3452, 0.1593, 0.5284, 0.4186, 0.3498, 0.267...

```
[41]: f, axe = plt.subplots(1,1, figsize=(18,6))

trained_models.sort_values(by=['Mean Cross Validated R2 Score'],
    ↪ascending=False, inplace=True)

sns.barplot(x='Mean Cross Validated R2 Score', y='Model', data =
    ↪trained_models, ax = axe)
axe.set_xlabel('Mean Cross Validated R2 Score', size=16)
axe.set_ylabel('Model')
axe.set_xlim(0,1.0)

plt.show()
```




```
[42]: f, axe = plt.subplots(1,1, figsize=(18,6))

trained_models.sort_values(by=['RMSE'], ascending=True, inplace=True)

sns.barplot(x='RMSE', y='Model', data = trained_models, ax = axe)
axe.set_xlabel('Root Mean Squared Error', size=16)
axe.set_ylabel('Model')
axe.set_xlim(0,1.0)

plt.show()
```



2 Klasyfikacja cen

Dzielimy zbiór danych na 5 przedziałów, aby móc przewidzieć etykiety zamiast wartości ciągłych.

```
[43]: import pandas as pd
X = cleaned_data.drop(columns='PRICE')
y = cleaned_data['PRICE']
y = np.round(y, decimals=-3) #przy cenach posiadłości część nie ma aż tak dużo
    ↪znaczenia znaczenia. Zakładając że ostatnie 2 miejsca wynoszą 99 to nawet
    ↪dla minimalnej wartości w zbiorze danych stonowi to jedynie 0.2%
```

```

sorted_prices = y.sort_values()

categories = pd.qcut(sorted_prices, q=3, labels=[ 'Low', 'Medium', 'High'])

category_ranges = {}
for category in categories.cat.categories:
    min_val = sorted_prices[categories == category].min()
    max_val = sorted_prices[categories == category].max()
    category_ranges[category] = (min_val, max_val)

print("Wartości dla każdego przedziału:")
for category, (min_val, max_val) in category_ranges.items():
    print(f"{category}: {min_val} - {max_val}")

```

Wartości dla każdego przedziału:

Low: 50000 - 599000

Medium: 600000 - 1175000

High: 1180000 - 6000000

```

[44]: y = y.apply(lambda x: 0 if x < 599000 else (1 if x < 1100000 else (2 if x < 2500000 else 3)))

print(y.value_counts())

```

PRICE

0 1477

1 1439

2 999

3 591

Name: count, dtype: int64

```

[45]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    random_state=42)

trained_models = []

```

2.1 Modele Klasyfikacji

2.1.1 DecisionTreeClassifier

```

[46]: from sklearn.metrics import confusion_matrix, classification_report
from sklearn.tree import DecisionTreeClassifier
clf_gini = DecisionTreeClassifier(criterion='gini', random_state=0)
clf_gini.fit(X_train, y_train)
y_pred_gini = clf_gini.predict(X_test)

conf_matrix = confusion_matrix(y_test, y_pred_gini)

```

```

class_report = classification_report(y_test, y_pred_gini)

print("Confusion Matrix:")
print(conf_matrix)
print("\nClassification Report:")
print(class_report)

```

Confusion Matrix:

```

[[325  99   9   3]
 [ 81 266  92   5]
 [ 24  87 151  32]
 [   2   7  36 133]]

```

Classification Report:

	precision	recall	f1-score	support
0	0.75	0.75	0.75	436
1	0.58	0.60	0.59	444
2	0.52	0.51	0.52	294
3	0.77	0.75	0.76	178
accuracy			0.65	1352
macro avg	0.66	0.65	0.65	1352
weighted avg	0.65	0.65	0.65	1352

2.1.2 Forest Classifier

```

[47]: from sklearn.ensemble import RandomForestClassifier
      from sklearn.metrics import confusion_matrix, classification_report

      # Instantiate the Random Forest Classifier
      clf_forest = RandomForestClassifier(random_state=0)
      clf_forest.fit(X_train, y_train)
      y_pred_forest = clf_forest.predict(X_test)
      conf_matrix_forest = confusion_matrix(y_test, y_pred_forest)
      class_report_forest = classification_report(y_test, y_pred_forest)

      print("Confusion Matrix (Random Forest Classifier):")
      print(conf_matrix_forest)
      print("\nClassification Report (Random Forest Classifier):")
      print(class_report_forest)

```

Confusion Matrix (Random Forest Classifier):

```

[[341  89   4   2]
 [ 67 317  56   4]
 [ 17  68 179  30]
 [   2   2  35 139]]

```

Classification Report (Random Forest Classifier):

	precision	recall	f1-score	support
0	0.80	0.78	0.79	436
1	0.67	0.71	0.69	444
2	0.65	0.61	0.63	294
3	0.79	0.78	0.79	178
accuracy			0.72	1352
macro avg	0.73	0.72	0.72	1352
weighted avg	0.72	0.72	0.72	1352

2.2 Podsumowanie Klasyfikacji

Ponieważ `roc_curve` i `roc_auc_score` nie obsługują oceny modeli do przewidywania wielu etykiet, postanowiliśmy porównać modele pod względem jakości przewidywania dla każdej klasy osobno.

```
[48]: import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, roc_auc_score
from sklearn.preprocessing import label_binarize

y_test_binarized = label_binarize(y_test, classes=np.unique(y_test))
n_classes = y_test_binarized.shape[1]

clf_gini = DecisionTreeClassifier(criterion='gini', random_state=0)
clf_gini.fit(X_train, y_train)
y_pred_gini_proba = clf_gini.predict_proba(X_test)

clf_forest = RandomForestClassifier(random_state=0)
clf_forest.fit(X_train, y_train)
y_pred_forest_proba = clf_forest.predict_proba(X_test)

fpr_gini = dict()
tpr_gini = dict()
roc_auc_gini = dict()
fpr_forest = dict()
tpr_forest = dict()
roc_auc_forest = dict()

for i in range(n_classes):
    fpr_gini[i], tpr_gini[i], _ = roc_curve(y_test_binarized[:, i],
    ↪ y_pred_gini_proba[:, i])
    roc_auc_gini[i] = roc_auc_score(y_test_binarized[:, i], y_pred_gini_proba[
    ↪ :, i])
```

```

    fpr_forest[i], tpr_forest[i], _ = roc_curve(y_test_binarized[:, i],
↪y_pred_forest_proba[:, i])
    roc_auc_forest[i] = roc_auc_score(y_test_binarized[:, i],
↪y_pred_forest_proba[:, i])

plt.figure(figsize=(12, 8))
colors = plt.cm.get_cmap('tab10', n_classes)

for i in range(n_classes):
    plt.plot(fpr_gini[i], tpr_gini[i], color=colors(i), lw=2,
             label=f'Decision Tree Class {i} (AUC = {roc_auc_gini[i]:.2f})')
    plt.plot(fpr_forest[i], tpr_forest[i], color=colors(i), linestyle='--',
↪lw=2,
             label=f'Random Forest Class {i} (AUC = {roc_auc_forest[i]:.2f})')

plt.plot([0, 1], [0, 1], 'k--', lw=2, label='Random Guess')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Multiclass Classification')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()


```

<ipython-input-48-2d643a2bc3e9>:31: MatplotlibDeprecationWarning: The get_cmap function was deprecated in Matplotlib 3.7 and will be removed two minor releases later. Use ``matplotlib.colormaps[name]`` or ``matplotlib.colormaps.get_cmap(obj)`` instead.

```

    colors = plt.cm.get_cmap('tab10', n_classes)

```



Z wykresu widać, że Forest Classifier jest zdecydowanie lepszym modelem dla przewidywania każdej z etykiet.

3 Wnioski

Ze względu na dobre dopasowanie, wybrałyśmy XGBRegressor oraz Forest Regressor Model (30 drzew) do dalszego etapu projektu. Spośród wielu wytrenowanych modeli, te dwa charakteryzowały się najlepszym dopasowaniem pod względem uśrednionej wartości współczynnika determinacji. Jeśli wyniki po przeprowadzeniu hiperparametryzacji będą niezadowalające, spróbujemy dopasować model klasyfikacji Forest Classifier.

hiperparametryzacja

June 9, 2024

1 Regresja

```
[ ]: from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from xgboost import XGBRegressor
import pandas as pd
from sklearn.ensemble import RandomForestRegressor

[ ]: cleaned_data = pd.read_excel('clean_data_relevant.xlsx')

[ ]: cleaned_data = cleaned_data[['BROKER', 'LATITUDE', 'LONGITUDE', 'TYPE',
    ↪ 'BOROUGH', 'NEIGHBOURHOOD', 'BEDS', 'POSTCODE', 'BATH', 'PROPERTYSQFT',
    ↪ 'PRICE']]

[ ]: from sklearn.model_selection import train_test_split

X = cleaned_data.drop(columns='PRICE')
y = cleaned_data['PRICE']

[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    ↪ random_state=42)
trained_models = []
```

1.1 Modele

```
[ ]: linear = LinearRegression()

param_grid = {
    'fit_intercept': [True, False],
    'copy_X': [True, False],
    'n_jobs': [3, 4, 10, 15, 20, 50, 80],
    'positive': [True, False]
}
```



```

grid_search = GridSearchCV(estimator=linear, param_grid=param_grid, cv=5,
    ↳scoring='neg_mean_squared_error', verbose=2)
grid_search.fit(X_train, y_train)
best_params = grid_search.best_params_
print("Best parameters found: ", best_params)

best_xgb_model = grid_search.best_estimator_
y_pred = best_xgb_model.predict(X_test)
r2 = r2_score(y_test, y_pred)
print("R-squared: ", r2)

```

Fitting 5 folds for each of 56 candidates, totalling 280 fits

```

[CV] END copy_X=True, fit_intercept=True, n_jobs=3, positive=True; total time=
0.0s
[CV] END copy_X=True, fit_intercept=True, n_jobs=3, positive=True; total time=
0.0s
[CV] END copy_X=True, fit_intercept=True, n_jobs=3, positive=True; total time=
0.0s
[CV] END copy_X=True, fit_intercept=True, n_jobs=3, positive=True; total time=
0.0s
[CV] END copy_X=True, fit_intercept=True, n_jobs=3, positive=True; total time=
0.0s
[CV] END copy_X=True, fit_intercept=True, n_jobs=3, positive=False; total time=
0.0s
[CV] END copy_X=True, fit_intercept=True, n_jobs=3, positive=False; total time=
0.0s
[CV] END copy_X=True, fit_intercept=True, n_jobs=3, positive=False; total time=
0.0s
[CV] END copy_X=True, fit_intercept=True, n_jobs=3, positive=False; total time=
0.0s
[CV] END copy_X=True, fit_intercept=True, n_jobs=3, positive=False; total time=
0.0s
[CV] END copy_X=True, fit_intercept=True, n_jobs=4, positive=True; total time=
0.0s
[CV] END copy_X=True, fit_intercept=True, n_jobs=4, positive=True; total time=
0.0s
[CV] END copy_X=True, fit_intercept=True, n_jobs=4, positive=True; total time=
0.0s
[CV] END copy_X=True, fit_intercept=True, n_jobs=4, positive=True; total time=
0.0s
[CV] END copy_X=True, fit_intercept=True, n_jobs=4, positive=False; total time=
0.0s
[CV] END copy_X=True, fit_intercept=True, n_jobs=4, positive=False; total time=
0.1s
[CV] END copy_X=True, fit_intercept=True, n_jobs=4, positive=False; total time=

```


[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

```

time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=10, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=10, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=10, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=10, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=10, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=15, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=15, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=15, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=15, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=15, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=15, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=15, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=15, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=20, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=20, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=20, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=20, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=20, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=20, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=20, positive=False; total

```

```

time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=20, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=20, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=50, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=50, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=50, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=50, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=50, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=50, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=50, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=50, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=50, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=80, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=80, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=80, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=80, positive=True; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=80, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=80, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=80, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=80, positive=False; total
time= 0.0s
[CV] END copy_X=False, fit_intercept=False, n_jobs=80, positive=False; total
time= 0.0s
Best parameters found: {'copy_X': True, 'fit_intercept': True, 'n_jobs': 3,
'positive': False}
R-squared: 0.35106361164412225

```

```
[ ]: xgb_model = XGBRegressor()

param_grid = {
    'n_estimators': [25, 50, 75]
}

grid_search = GridSearchCV(estimator=xgb_model, param_grid=param_grid, cv=5,
    ↪scoring='neg_mean_squared_error')
grid_search.fit(X_train, y_train)

best_params = grid_search.best_params_
print("Best parameters found: ", best_params)

best_xgb_model = grid_search.best_estimator_
y_pred = best_xgb_model.predict(X_test)
r2 = r2_score(y_test, y_pred)
print("R-squared: ", r2)
```

Best parameters found: {'n_estimators': 50}
R-squared: 0.6301618111670517

```
[ ]: xgb_model = XGBRegressor()

param_grid = {
    'n_estimators': [25, 50, 75],
    'max_depth': [ 3, 5, 10]
}

grid_search = GridSearchCV(estimator=xgb_model, param_grid=param_grid, cv=5,
    ↪scoring='neg_mean_squared_error')
grid_search.fit(X_train, y_train)
best_params = grid_search.best_params_
print("Best parameters found: ", best_params)

best_xgb_model_2 = grid_search.best_estimator_
y_pred = best_xgb_model_2.predict(X_test)
r2 = r2_score(y_test, y_pred)
print("R-squared: ", r2)
```

Best parameters found: {'max_depth': 3, 'n_estimators': 25}
R-squared: 0.6634425443734272

```
[ ]: xgb_model = XGBRegressor()

param_grid = {
    'n_estimators': [25, 50, 75],
    'max_depth': [ 3, 5, 10],
```

```

        'learning_rate': [0.05, 0.1, 0.15]
    }

    grid_search = GridSearchCV(estimator=xgb_model, param_grid=param_grid, cv=5,
                               scoring='neg_mean_squared_error')
    grid_search.fit(X_train, y_train)
    best_params = grid_search.best_params_
    print("Best parameters found: ", best_params)

    best_xgb_model_3 = grid_search.best_estimator_
    y_pred = best_xgb_model_3.predict(X_test)
    r2 = r2_score(y_test, y_pred)
    print("R-squared: ", r2)

```

Best parameters found: {'learning_rate': 0.15, 'max_depth': 3, 'n_estimators': 50}

R-squared: 0.6532830889068985

```

[ ]: xgb_model = XGBRegressor()

    param_grid = {
        'n_estimators': [25, 50, 75],
        'max_depth': [3, 5, 10],
        'learning_rate': [0.1, 0.15],
        'max_leaves': [25, 50, 200, 0]
    }

    grid_search = GridSearchCV(estimator=xgb_model, param_grid=param_grid, cv=5,
                               scoring='neg_mean_squared_error')
    grid_search.fit(X_train, y_train)

    best_params = grid_search.best_params_
    print("Best parameters found: ", best_params)

    best_xgb_model_4 = grid_search.best_estimator_
    y_pred = best_xgb_model_4.predict(X_test)
    r2 = r2_score(y_test, y_pred)
    print("R-squared: ", r2)

```

Best parameters found: {'learning_rate': 0.15, 'max_depth': 3, 'max_leaves': 25, 'n_estimators': 50}

R-squared: 0.6532830889068985

```

[ ]: xgb_model = XGBRegressor()

    param_grid = {

```

```

        'n_estimators': [25, 50, 75],
        'max_depth': [ 3, 5, 10],
        'learning_rate': [0.1, 0.15],
        'max_leaves': [25, 50, 200, 0],
        'tree_method': ['exact', 'approx', 'hist']
    }

    grid_search = GridSearchCV(estimator=xgb_model, param_grid=param_grid, cv=5,
                               ↪scoring='neg_mean_squared_error')
    grid_search.fit(X_train, y_train)

    best_params = grid_search.best_params_
    print("Best parameters found: ", best_params)

    best_xgb_model_5 = grid_search.best_estimator_
    y_pred = best_xgb_model_5.predict(X_test)
    r2 = r2_score(y_test, y_pred)
    print("R-squared: ", r2)

```

Best parameters found: {'learning_rate': 0.15, 'max_depth': 3, 'max_leaves': 25, 'n_estimators': 75, 'tree_method': 'exact'}

R-squared: 0.6414683850465416

```

[ ]: xgb_model = XGBRegressor()

    param_grid = {
        'n_estimators': [25, 50, 75],
        'max_depth': [ 5, 8, 10],
        'learning_rate': [0.05, 0.1, 0.15],
        'max_leaves': [25, 50, 200, 0],
        'tree_method': ['exact', 'approx', 'hist'],
        'gamma': [0, 0.1, 0.2]
    }

    grid_search = GridSearchCV(estimator=xgb_model, param_grid=param_grid, cv=5,
                               ↪scoring='neg_mean_squared_error')
    grid_search.fit(X_train, y_train)
    best_params = grid_search.best_params_
    print("Best parameters found: ", best_params)

    best_xgb_model_6 = grid_search.best_estimator_
    y_pred = best_xgb_model_6.predict(X_test)
    r2 = r2_score(y_test, y_pred)
    print("R-squared: ", r2)

```

Best parameters found: {'gamma': 0, 'learning_rate': 0.1, 'max_depth': 5,

```
'max_leaves': 25, 'n_estimators': 75, 'tree_method': 'exact'}
R-squared: 0.6741242999093524
```

```
[ ]: xgb_model = XGBRegressor()

param_grid = {
    'n_estimators': [25, 50, 75],
    'max_depth': [ 5, 8, 10],
    'learning_rate': [0.05, 0.1, 0.15],
    'max_leaves': [25, 50, 200, 0],
    'tree_method': ['exact', 'approx', 'hist'],
    'gamma': [0.1, 0.2]
}

grid_search = RandomizedSearchCV(estimator=xgb_model,
    ↪ param_distributions=param_grid, cv=5, scoring='neg_mean_squared_error',
    ↪ n_iter=200)
grid_search.fit(X_train, y_train)
best_params = grid_search.best_params_
print("Best parameters found: ", best_params)

best_xgb_model_7 = grid_search.best_estimator_
y_pred = best_xgb_model_7.predict(X_test)
r2 = r2_score(y_test, y_pred)
print("R-squared: ", r2)
```

```
Best parameters found: {'tree_method': 'exact', 'n_estimators': 75,
'max_leaves': 25, 'max_depth': 5, 'learning_rate': 0.1, 'gamma': 0.2}
R-squared: 0.6741242999093524
```

```
[ ]: xgb_model = XGBRegressor()

param_grid = {
    'n_estimators': [100, 500, 750],
    'max_depth': [ 5, 7],
    'learning_rate': [0.05, 0.1],
    'max_leaves': [0, 10],
    'tree_method': [ 'approx', 'exact']
}

grid_search = RandomizedSearchCV(estimator=xgb_model,
    ↪ param_distributions=param_grid, cv=5, scoring='neg_mean_squared_error',
    ↪ n_iter=200)
grid_search.fit(X_train, y_train)
best_params = grid_search.best_params_
print("Best parameters found: ", best_params)
```



```

best_xgb_model_8 = grid_search.best_estimator_
y_pred = best_xgb_model_8.predict(X_test)
r2 = r2_score(y_test, y_pred)
print("R-squared: ", r2)

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/model_selection/_search.py:305:
UserWarning: The total space of parameters 48 is smaller than n_iter=200.
Running 48 iterations. For exhaustive searches, use GridSearchCV.
  warnings.warn(

Best parameters found: {'tree_method': 'exact', 'n_estimators': 500,
'max_leaves': 0, 'max_depth': 5, 'learning_rate': 0.05}
R-squared: 0.6813970124277131

```

```

[ ]: rf_model = RandomForestRegressor()

param_grid = {
    'n_estimators': [20,50, 70],
    'max_depth': [None, 10, 20, 30],
    'criterion':['squared_error', 'absolute_error', 'friedman_mse', 'poisson']
}

grid_search = GridSearchCV(estimator=rf_model, param_grid=param_grid, cv=5,
    ↪scoring='neg_mean_squared_error')
grid_search.fit(X_train, y_train)
best_params = grid_search.best_params_
print("Best parameters found: ", best_params)

best_rf_model = grid_search.best_estimator_
y_pred = best_rf_model.predict(X_test)
r2 = r2_score(y_test, y_pred)
print("R-squared: ", r2)

```

```

Best parameters found: {'criterion': 'poisson', 'max_depth': 20,
'n_estimators': 50}
R-squared: 0.6630910611411529

```

```

[ ]: rf_model = RandomForestRegressor()

param_grid = {
    'n_estimators': [20,50, 70],
    'max_features': ['auto', 'sqrt', 'log2', None],
    'bootstrap':[True , False]
}

grid_search = GridSearchCV(estimator=rf_model, param_grid=param_grid, cv=5,
    ↪scoring='neg_mean_squared_error')
grid_search.fit(X_train, y_train)

```

```

best_params = grid_search.best_params_
print("Best parameters found: ", best_params)

best_rf_model_2 = grid_search.best_estimator_
y_pred = best_rf_model_2.predict(X_test)
r2 = r2_score(y_test, y_pred)
print("R-squared: ", r2)

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or
remove this parameter as it is also the default value for RandomForestRegressors
and ExtraTreesRegressors.

```

```

warn(

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or
remove this parameter as it is also the default value for RandomForestRegressors
and ExtraTreesRegressors.

```

```

warn(

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or
remove this parameter as it is also the default value for RandomForestRegressors
and ExtraTreesRegressors.

```

```

warn(

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or
remove this parameter as it is also the default value for RandomForestRegressors
and ExtraTreesRegressors.

```

```

warn(

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or
remove this parameter as it is also the default value for RandomForestRegressors
and ExtraTreesRegressors.

```

```

warn(

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or
remove this parameter as it is also the default value for RandomForestRegressors
and ExtraTreesRegressors.

```

```

warn(

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or

```

remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or
```

remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or
```

remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
  
Best parameters found: {'bootstrap': True, 'max_features': 'sqrt',  
'n_estimators': 70}
```

R-squared: 0.6682914367280991

```
[ ]: rf_model = RandomForestRegressor()

param_grid = {
    'n_estimators': [20, 50],
    'max_depth': [None, 5, 10],
    'criterion': ['squared_error', 'friedman_mse', 'poisson'],
    'max_features': ['auto', 'sqrt', 'log2', None],
    'bootstrap': [True, False]
}

grid_search = GridSearchCV(estimator=rf_model, param_grid=param_grid, cv=5,
    ↪scoring='neg_mean_squared_error')
grid_search.fit(X_train, y_train)
best_params = grid_search.best_params_
print("Best parameters found: ", best_params)

best_rf_model_3 = grid_search.best_estimator_
y_pred = best_rf_model_3.predict(X_test)
r2 = r2_score(y_test, y_pred)
print("R-squared: ", r2)
```

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or
remove this parameter as it is also the default value for RandomForestRegressors
and ExtraTreesRegressors.

warn(

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or
remove this parameter as it is also the default value for RandomForestRegressors
and ExtraTreesRegressors.

warn(

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or
remove this parameter as it is also the default value for RandomForestRegressors
and ExtraTreesRegressors.

warn(

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or
remove this parameter as it is also the default value for RandomForestRegressors
and ExtraTreesRegressors.

warn(

/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```


FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```


FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```


FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:
```

FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or remove this parameter as it is also the default value for RandomForestRegressors and ExtraTreesRegressors.

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_forest.py:413:  
FutureWarning: `max_features='auto'` has been deprecated in 1.1 and will be  
removed in 1.3. To keep the past behaviour, explicitly set `max_features=1.0` or  
remove this parameter as it is also the default value for RandomForestRegressors  
and ExtraTreesRegressors.
```

```
warn(  

```

```
Best parameters found: {'bootstrap': True, 'criterion': 'poisson', 'max_depth': 10, 'max_features': 'sqrt', 'n_estimators': 50}
R-squared: 0.6675360383586544
```

Po hiperparametryzacji najlepszym modelem okazał się XGBRegressor o następujących hiperparametrach:

- 'gamma': 0,
- – 'learning_rate': 0.1,
- – 'max_depth': 5,
- – 'max_leaves': 25,
- – 'n_estimators': 75,
- 'tree_method': 'exact'


RY2 jest równe dla tego modelu 0.674, co wskazuje na umiarkowanie silny model regresji. Oznacza to, że 67.4% zmienności cen jest wyjaśniana przez zmienne niezależne użyte w szkoleniu modelu regresji.

Ważność cech

```
[ ]: import lime
import shap

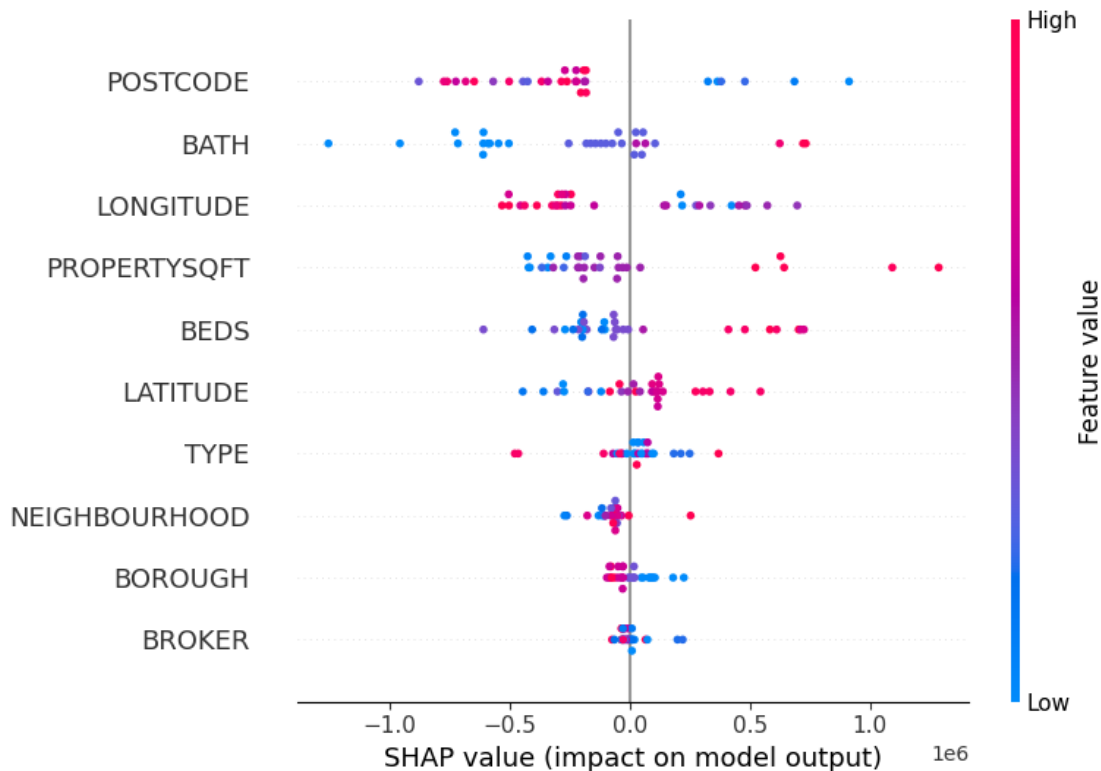
[ ]: feat_importances = pd.DataFrame(best_xgb_model_7.feature_importances_,
    ↪index=X_train.columns, columns=["Importance"])
feat_importances.sort_values(by='Importance', ascending=False, inplace=True)
feat_importances.plot(kind='bar', figsize=(8,6))
shap.initjs()
```

<IPython.core.display.HTML object>



Na podstawie powyższego wykresu można stwierdzić, że czynniki takie jak liczba łazienek, lokalizacja i wielkość nieruchomości mają kluczowe znaczenie dla określania cen mieszkań. Cechy o niższych wynikach ważności nadal wnoszą wkład do modelu, ale mają mniejszy wpływ na ostateczną prognozę.

```
[ ]: explainer = shap.Explainer(best_xgb_model_7, feature_names = X_train.columns)
shap_values = explainer.shap_values(X_test)
shap_values = explainer.shap_values(X_test[:30])
shap.summary_plot(shap_values, X_test[:30])
```



Rozrzut wartości SHAP dla każdej cechy pokazuje zakres jej wpływu na prognozy.



- POSTCODE: Wysokie wartości SHAP dla określonych kodów pocztowych sugerują, że lokalizacja może znacząco zwiększyć lub zmniejszyć przewidywaną cenę mieszkania. Podobne wnioski można wyciągnąć dla zmiennych LONGITUDE i LATITUDE.
- BATH, BEDS: Liczba łazienek (sypialni) również ma znaczący wpływ - większa liczba łazienek (sypialni) zwiększa przewidywaną cenę.
- PROPERTYSQFT: Większe rozmiary nieruchomości powodują wyższe ceny nieruchomości.
- TYPE, BOROUGH, NEIGHBOURHOOD, BROKER: Cechy mają mniejszy, ale nadal zauważalny wpływ na ceny mieszkań.

```
[ ]: explainer2 = lime.lime_tabular.LimeTabularExplainer(X_train.values,
    ↪ feature_names=X_train.columns.values.tolist(), class_names=['PRICE'],
    ↪ verbose=True, mode='regression')

instance = X_train.iloc[4].values.reshape(1, -1)
explanation2 = explainer2.explain_instance(X_test.values[1000],
    ↪ best_xgb_model_7.predict, num_features=8)
explanation2.as_pyplot_figure()
```

```
Intercept 3499325.1080469214
Prediction_local [374777.08142794]
Right: 262017.66
```

[]:



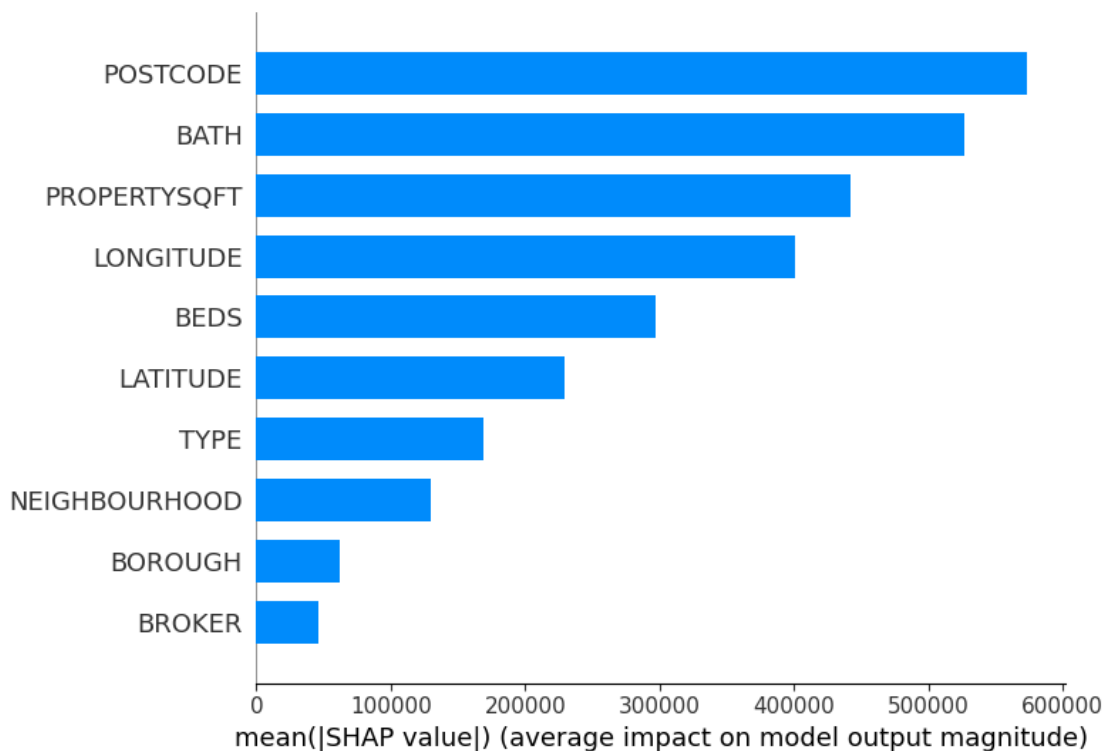
BATH ≤ 0.10 : Liczba łazienek znacząco wpływa na obniżenie prognozy ceny gdy przeskalowana wartość jest mniejsza lub równa od 0.10. 0.11 < PROPERTYSQFT ≤ 0.23 : oznacza, że jeżeli przeskalowana wartość powierzchni nieruchomości mieści się w przedziale od 0,11 do 0,23, to ma

ona znaczny negatywny wpływ na prognozowaną cenę $POSTCODE > 128.00$: gdy wartość kategorii przypisana kodowi pocztowemu jest większa niż 128.00 to przewidywana cena jest niższa. $LONGITUDE > 0.70$: gdy przeskalowana wartość długości geograficznej jest większa niż 0.70, to prognozowana cena nieruchomości jest niższa. $TYPE \leq 0.00$: Gdy typ budynku jest 'co-op' to zwiększa to cenę nieruchomości. $BEDS \leq 0.17$: Liczba łazienek wpływa na obniżenie prognozy ceny gdy jej przeskalowana wartość jest mniejsza lub równa od 0.17. $0.55 < LATITUDE \leq 0.66$: Szerokość geograficzna pozytywnie wpływa na cenę nieruchomości gdy jej przeskalowana wartość mieści się w przedziale od 0.55 do 0.66. $NEIGHBOURHOOD \leq 91.00$: To oznacza, że gdy wartość przyporządkowanej kategorii sąsiedztwa jest mniejsza lub równa 91.00, to prognozowana cena nieruchomości jest niższa - wpływ ten jest zdecydowanie mniejszy w porównaniu do innych cech.

```
[ ]: shap_values = explainer.shap_values(X_test)
      print("Variable Importance Plot - Global Interpretation")

      shap.summary_plot(shap_values, X_test, plot_type="bar")
```

Variable Importance Plot - Global Interpretation



Wykres przedstawia globalną interpretację ważności zmiennych modelu za pomocą wartości SHAP. Najbardziej wpływowymi cechami są zmienne: POSTCODE, BATH i PROPERTYSQFT. Pozostałymi mniej istotnymi cechami są LONGITUDE, BEDS, LATITUDE, TYPE, NEIGHBOURHOOD, BOROUGH i BROKER, które również przyczyniają się do dokładności modelu.

Ze względu na umiarkowanie dobrą jakość wytrenowanych modeli regresyjnych, zdecydowaliśmy się nie skupiać na klasyfikacji.